ICE 2023 Pre-Proceedings

Clément Aubert   Cinzia Di Giusto   Simon Fowler   Larisa Safina

30th May 2023

This document contains the *informal* pre-proceedings of the 16th Interaction and Concurrency Experience (ICE 2023). The post-proceedings will be published on EPTCS.

Contents

Research Challenges in Orchestration Synthesis

Davide Basile   Maurice H. ter Beek

Formal Methods and Tools lab, ISTI–CNR, Pisa, Italy

{davide.basile,maurice.terbeek}@isti.cnr.it

Contract automata allow to formally define the behaviour of service contracts in terms of service offers and requests, some of which are moreover optional and some of which are necessary. A composition of contracts is said to be in agreement if all service requests are matched by corresponding offers. Whenever a composition of contracts is not in agreement, it can be refined to reach an agreement using the orchestration synthesis algorithm. This algorithm is a variant of the synthesis algorithm used in supervisory control theory and it is based on the fact that optional transitions are controllable, whereas necessary transitions are at most semi-controllable and cannot always be controlled. In fact, the resulting orchestration is such that as much of the behaviour in agreement is maintained. In this paper, we discuss recent developments of the orchestration synthesis algorithm for contract automata. Notably, we present a refined notion of semi-controllability and compare it with the original notion by means of examples. We then discuss the current limits of the orchestration synthesis algorithm and identify a number of research challenges together with a research roadmap.

1 Introduction

Orchestrations of services describe how control and data exchanges are coordinated in distributed service-based applications and systems. Their principled design is identified in [16] as one of the primary research challenges for the next 10 years, and the Service Computing Manifesto [16] points out that "Service systems have so far been built without an adequate rigorous foundation that would enable reasoning about them" and, moreover, that "The design of service systems should build upon a formal model of services".

The problem of synthesising well-behaving orchestrations of services can be viewed as a specific instance of the more general problem of synthesising strategies in games [9, 7]. This can be solved using refined algorithms from supervisory control for discrete event systems [24, 1], which have well-established relationships with reactive systems synthesis [20], parity games [23], automated behaviour composition [21] and automated planning [17].

Contract automata are a specific type of finite state automata that are used to formally define the behaviour of service contracts. These automata express contracts in terms of both offers and requests [10]. When multiple contracts are composed, they are said to be in agreement if all service requests from one contract are matched by another contract's corresponding offers. A composition of contracts that is not in agreement, can automatically be refined to reach an agreement by means of the orchestration synthesis algorithm, which is a variation of the synthesis algorithm used in supervisory control theory. This orchestration synthesis algorithm for contract automata is described in [8, 9].

The classic algorithm for synthesising a most permissive controller distinguishes transitions whose controllability is invariant [24, 1]. In service contracts, instead, the controllability of certain transitions may vary depending on specific conditions on the orchestration of contracts [9]. The contract automata library `CATLib` [5] implements contract automata and their operations (e.g., composition and synthesis). Orchestrations of contract automata abstract from their underlying realisation; an orchestrator is assumed

to interact with the services to realise the overall behaviour as prescribed by the orchestration contract. The contract automata runtime environment CARE [6] implements an orchestrator that interprets the synthesised orchestration to coordinate the services, where each service is implementing a contract. Thus, CARE is explicating the low-level interactions that are abstracted in contract automata orchestrations. Notably, one aspect that is abstracted in contract automata and concretised at the implementation level is that of selecting the next transition to execute in the presence of choice. In [6], different implementations are proposed based on whether services may participate externally or internally in a choice.

This paper delves into challenges and research issues for orchestration synthesis of contract automata, given the latest developments in this field. In particular, we start by refining the current definition of semi-controllability to consider the aforementioned possible realisations of choices defined in [6]. We provide several examples to illustrate the differences between the refined definition and the original definition. The various definitions of semi-controllability lead to different sets of contract automata orchestrations, which we present in Figure 3 together with an example for each level of the orchestration hierarchy depicted. This allows us to highlight the unique characteristics of each level and to identify current issues in synthesising orchestrations of contract automata using these examples. Based on the issues presented, we then outline future research challenges in the orchestration synthesis of contract automata and a research roadmap to address them.

Related Work   At last year's ICE 2022 workshop, the compositionality of communicating finite state machines (CFSM) with asynchronous semantics was discussed in [3]. Also contract automata are composable, enabling the modelling of systems of systems. Moreover, under certain specific conditions that were presented at the 2014 edition of ICE [11, 12], an orchestration of contract automata can be translated into a choreography of synchronous or asynchronous CFSM. The relation between multiparty session types and CFSM is discussed in [27]. Therefore, contract automata can be related to multiparty session types by exploiting their common relation with CFSM [11, 12, 27].

The contract automata approach is closer to [22], in which behavioural types are expressed as finite state automata of Mungo, called typestates [25]. Similarly to CARE, the runtime environment for contract automata [6], in Mungo finite state automata are used as behaviour assigned to Java classes (one automaton per class), with transition labels corresponding to methods of the classes. A tool to translate typestates into automata was presented at ICE 2020 [26]. CATApp, a graphical front-end tool for designing contract automata, is available in [19]. A tool similar to Mungo is JaTyC (Java Typestate Checker) [2].

The refined definition of semi-controllability presented in this paper closely aligns with the notion of weak receptiveness in team automata [14, 15]. However, the challenges addressed in this paper are primarily related to the problem of synthesising an orchestration of services and as such are not directly relevant to team automata.

Differently from the semi-controllability for orchestrations, a distinct notion of semi-controllability has been studied in [9, 4] for choreographies of services. Finally, while a runtime environment for the orchestration of services has been proposed in [6], this has yet to be realised for the case of choreographies, which could result in improvements in the notion of semi-controllability for choreographies.

Outline   We start by providing some background on contract automata and orchestration synthesis in Section 2. We introduce a refined notion of semi-controllability in Section 3. In Section 4, we present several research challenges for orchestration synthesis of contract automata. We conclude in Section 5.

2   Background

We will begin by formally introducing contract automata and their synthesis operation. Contract automata are a type of finite state automata that use a partitioned alphabet of actions. A Contract Automaton (CA) can model either a single service or a composition of multiple services that perform actions. The number of services in a CA is known as its rank. If the rank of a CA is 1, then the contract is referred to as a principal (i.e., a single service).

The labels of a CA are vectors of atomic elements known as actions. Actions are categorised as either requests (prefixed by ?), offers (prefixed by !), or idle actions (represented by a distinguished symbol $-$). Requests and offers belong to the sets R and O, respectively, and they are pairwise disjoint. The states of a CA are vectors of atomic elements known as basic states. Labels are restricted to requests, offers or matches. In a request (resp. offer) label there is a single request (resp. offer) action and all other actions are idle. In a match label there is a single pair of request and offer actions that match, and all other actions are idle. The length of the vectors of states and labels is equal to the rank of the CA. For example, the label $[!a, ?a, -, -]$ is a match where the request action $?a$ is matched by the offer action $!a$, and all other actions are idle. Note the difference between a request label (e.g., $[?a, -]$) and a request action (e.g., $?a$). A transition may also be called a request, offer or match according to its label. Figure 4 depicts three principal contracts, whilst Figure 5 depicts a contract of rank 3.

The goal of each service is to reach an accepting (*final*) state such that all its request (and possibly offer) actions are matched. Transitions are equipped with *modalities*, i.e., *necessary* ($\square$) and *optional* ($\circ$) transitions, respectively [1]. Optional transitions are controllable, whereas necessary transitions can be uncontrollable (called *urgent* necessary transitions) or semi-controllable (called *lazy* necessary transitions). The resulting formalism is called *Modal Service Contract Automata* (MSCA). In the following definition, given a vector $\vec{a}$, its ith element is denoted by $\vec{a}_{(i)}$.

Definition 1 (MSCA)**.** *Given a finite set of states $\mathcal{Q} = \{q_1, q_2, \ldots\}$, an MSCA $\mathcal{A}$ of rank n is a tuple $\langle Q, \vec{q}_0, A^r, A^o, T, F \rangle$, with set of states $Q = Q_1 \times \ldots \times Q_n \subseteq \mathcal{Q}^n$, initial state $\vec{q}_0 \in Q$, set of requests $A^r \subseteq \mathsf{R}$, set of offers $A^o \subseteq \mathsf{O}$, set of final states $F \subseteq Q$, set of transitions $T \subseteq Q \times A \times Q$, where $A \subseteq (A^r \cup A^o \cup \{\bullet\})^n$, partitioned into* optional *transitions T° and* necessary *transitions $T^\square$, with $T^\square$ further partitioned into* urgent *necessary transitions $T^{\square_u}$ and* lazy *necessary transitions $T^{\square_l}$, and such that given $t = (\vec{q}, \vec{a}, \vec{q}\,') \in T$: i) $\vec{a}$ is either a request, an offer or a match; ii) if $\vec{a}$ is an offer, then $t \in T^\circ$; and iii) $\forall i \in 1 \ldots n$, $\vec{a}_{(i)} = \bullet$ implies $\vec{q}_{(i)} = \vec{q}\,'_{(i)}$.*

Composition of services is rendered through the composition of their MSCA models by means of the *composition operator* $\otimes$, which is a variant of a synchronous product. This operator basically interleaves or matches the transitions of the component MSCA, but, whenever two component MSCA are enabled to execute their respective request/offer action, then the match is forced to happen. Moreover, a match involving a necessary transition of an operand is itself necessary. The rank of the composed MSCA is the sum of the ranks of its operands. The vectors of states and actions of the composed MSCA are built from the vectors of states and actions of the component MSCA, respectively. In this paper, we will only consider principal contracts and compositions of principals, which will be automatically refined into orchestrations (as shown in Figure 2). However, it is important to note that contracts can be created by composing contracts with a rank of one or higher.

In a composition of MSCA, typically various properties are analysed. We are especially interested in *agreement*. The property of agreement requires to match all requests, whereas offers can go unmatched.

---

[1]Originally, in [8], the optional modality was called permitted and denoted with $\diamondsuit$. Since in contract automata the two modalities are a partition, the terminology has been updated to avoid confusion with modal transition systems, where $\square \subseteq \diamondsuit$.

CA support the synthesis of the most permissive controller (mpc) known from the theory of supervisory control of discrete event systems [24, 18], where a finite state automaton model of a *supervisory controller* is synthesised from given (component) finite state automata that are composed. The synthesised automaton, if successfully generated (i.e., non-empty), is such that it is *non-blocking*, *controllable*, and *maximally permissive*. An automaton is said to be *non-blocking* if, from each state, at least one of the *final states* (distinguished stable states that represent completed 'tasks' [24]) can be reached without passing through so-called *forbidden states*, meaning that there is always a possibility to return to an accepted stable state (e.g., a final state).

The synthesised automaton is said to be *controllable* when only controllable transitions are disabled. Indeed, the supervisory controller is not permitted to directly block uncontrollable transitions from occurring; the controller is only allowed to disable them by preventing controllable actions from occurring. Finally, the fact that the resulting supervisory controller is said to be *maximally permissive* (or least restrictive) means that as much behaviour of the uncontrolled system as possible is present in the controlled system without violating neither the requirements, nor controllability nor the non-blocking condition.

Orchestration Synthesis As stated previously, optional transitions are controllable, whereas necessary transitions can be either uncontrollable (called *urgent*) or semi-controllable (called *lazy*). In the mpc synthesis (implemented in CATLib [9, 5]), all necessary transitions are *urgent*, i.e., they are always uncontrollable. This stems from the fact that traditionally uncontrollable transitions relate to an unpredictable environment.

When synthesising an orchestration of services, all necessary transitions are instead *lazy*, i.e., they are *semi-controllable* [8, 9]. A semi-controllable transition t is a transition that is either uncontrollable or controllable according to given conditions. In [9], different conditions are given according to whether the synthesis of an orchestration or a choreography is computed. In this paper, we only consider orchestrations. Below, we denote with $Dangling(\mathcal{A})$ the set of states that are not reachable from the initial state or cannot reach any final state. More in detail, a semi-controllable transition t is controllable if in a given portion $\mathcal{A}'$ of $\mathcal{A}$ there exists a semi-controllable match transition t', with source and target states not dangling, such that in both t and t' the *same* service, in the *same* local state, does the *same* request. Otherwise, t is uncontrollable.

Definition 2 (Controllability). *Let $\mathcal{A}$ be an MSCA and let $t = (\vec{q}_1, \vec{a}_1, \vec{q}_1') \in T_{\mathcal{A}}$. Then:*

- *if $t \in T_{\mathcal{A}}^{\circ}$, then t is* controllable *(in $\mathcal{A}$);*

- *if $t \in T_{\mathcal{A}}^{\square_u}$, then t is* uncontrollable *(in $\mathcal{A}$);*

- *if $t \in T_{\mathcal{A}}^{\square_l}$, then t is* semi-controllable *(in $\mathcal{A}$).*

Moreover, given $\mathcal{A}' \subseteq \mathcal{A}$, if t is semi-controllable and $\exists t' = (\vec{q}_2, \vec{a}_2, \vec{q}_2') \in T_{\mathcal{A}'}^{\square}$ in $\mathcal{A}'$ such that $\vec{a}_2$ is a match, $\vec{q}_2, \vec{q}_2' \notin Dangling(\mathcal{A}')$, $\vec{q}_{1(i)} = \vec{q}_{2(i)}$, and $\vec{a}_{1(i)} = \vec{a}_{2(i)} = ?a$ for some $i \in 0 \ldots rank(\mathcal{A})$, then t is controllable *in $\mathcal{A}'$ (via t'). Otherwise, t is* uncontrollable *in $\mathcal{A}'$.*

The interpretation of optional/controllable and urgent/uncontrollable transitions is standard [24, 18]. In the upcoming section, we will delve into different understandings and interpretations of the concept of semi-controllability. We remark that the orchestration synthesis defined below does not support urgent transitions. The orchestration synthesis, as defined below, involves an iterative refinement of the initial automaton $\mathcal{A}$ (i.e., the composition of contracts). In each iteration, transitions are selectively pruned, and a set R of forbidden states is updated accordingly. A transition t is pruned under one of two conditions: if it is a request (thus violating the agreement property enforced by the orchestration), or if the target

Client1 **Client2** **Server** $\mathbf{O}(\textbf{Client1} \otimes \textbf{Client1})$

Figure 1: Contracts of **Client1**, **Client2** and **Server**, and orchestration $\mathbf{O}(\textbf{Client1} \otimes \textbf{Client1})$

$\mathbf{O}(\textbf{Server} \otimes \textbf{Client2} \otimes \textbf{Client2})$

Figure 2: Orchestration $\mathbf{O}(\textbf{Server} \otimes \textbf{Client2} \otimes \textbf{Client2})$

state of t belongs to the set R computed up to that point. During the first iteration, all request transitions, including both lazy and optional ones, are pruned.

In Definition 2, the automaton $\mathcal{A}'$ represents an intermediate refinement of $\mathcal{A}$ (the starting composition) which occurs during an iteration of the synthesis process. Intuitively, the semi-controllable transition t of $\mathcal{A}$ is controllable in $\mathcal{A}'$ because there is another transition t' in $\mathcal{A}'$ matching the same request from the same service in the same state. Otherwise, if there is no such transition t' in $\mathcal{A}'$, then t is uncontrollable. Put differently, the controllability of t in $\mathcal{A}'$ relies on the presence of a corresponding transition t' within $\mathcal{A}'$ itself. If such a matching transition t' does not exist in $\mathcal{A}'$, then t is deemed uncontrollable.

Note that in Definition 2, it is not required for t and t' to be distinct. This implies that during the synthesis process, a semi-controllable match transition t can switch from being controllable to uncontrollable only after it has been pruned in a previous iteration. To clarify further, a semi-controllable match transition t can switch its controllability status from controllable to uncontrollable only when t is absent in the sub-automaton $\mathcal{A}'$ during the current iteration. If t is present in $\mathcal{A}'$ (i.e., it has not been pruned thus far), then, according to Definition 2, t is considered semi-controllable and controllable within $\mathcal{A}'$ via t itself. It is important to note that these considerations are applicable only if t is a match. Additionally, it is never the case that a semi-controllable transition t switches from uncontrollable to controllable since transitions are only removed during the synthesis process and are never added back.

The set R of forbidden states is updated at each iteration by adding source states of uncontrollable transitions and dangling states of the refined automaton in the current iteration. Specifically, when the synthesis process eliminates all transitions t' that satisfy the conditions for rendering the semi-controllable transition t controllable via t', then t becomes uncontrollable within the sub-automaton in the current iteration. It is worth noting that even if t was previously pruned in an earlier iteration, its source state $\vec{q}_1$ might still be reachable in the sub-automaton of the current iteration. Consequently, $\vec{q}_1$ is added

to the set R. In the subsequent iteration, all transitions with target state $\vec{q}_1$ will be pruned. This pruning of transitions whose target is $\vec{q}_1$ can potentially render another previously pruned semi-controllable transition as uncontrollable, thereby adding its source state to the updated set R. This refinement process continues until no further transitions are pruned, and no additional states are added to R. The resulting refined automaton obtained at the end of the synthesis process represents the orchestration automaton.

The algorithm for synthesising an orchestration enforcing agreement of MSCA is defined below.

Definition 3 (MSCA orchestration synthesis). *Let $\mathcal{A}$ be an MSCA and let $\mathcal{K}_0 = \mathcal{A}$ and $R_0 = Dangling(\mathcal{K}_0)$. We let the* orchestration synthesis function $f_o : MSCA \times 2^Q \rightarrow MSCA \times 2^Q$ *be defined as follows:*

$$f_o(\mathcal{K}_{i-1}, R_{i-1}) = (\mathcal{K}_i, R_i), \text{ with}$$
$$T_{\mathcal{K}_i} = T_{\mathcal{K}_{i-1}} \smallsetminus \{(\vec{q} \rightarrow \vec{q}') = t \in T_{\mathcal{K}_{i-1}} \mid (\vec{q}' \in R_{i-1} \vee t \text{ is a request})\}$$
$$R_i = R_{i-1} \cup \{\vec{q} \mid (\vec{q} \rightarrow) \in T_{\mathcal{A}}^{\square_l} \text{ is uncontrollable in } \mathcal{K}_i\} \cup Dangling(\mathcal{K}_i)$$

The orchestration automaton is obtained from the fixpoint of the function f_o. In the rest of the paper, if not stated otherwise, all necessary transitions in the examples are lazy (cf. Definition 1); for brevity and less cluttering in the figures, we denote them by $\square$ rather than $\square_l$.

Example 1. We provide an illustrative example to underline the differences between optional transitions, urgent necessary transitions and lazy necessary transitions. Figure 1 shows two client contracts and a server contract. Firstly, we discuss the difference between optional and necessary transitions. When all actions of the client contract are optional (**Client1**), there exists an orchestration of the composition of two **Client1** contracts, also depicted in Figure 1 (**O**(**Client1** ⊗ **Client1**)). Indeed the (transition labelled with the) request $?a$ is optional and can be removed to obtain the orchestration. If instead the request $?a$ was necessary (**Client2**), then there would be no orchestration for the composition of two **Client2** contracts, because the necessary request is never matched by a corresponding offer.

To illustrate the distinction between urgent and lazy necessary transitions, we consider also the server contract shown in Figure 1. If we were to employ the traditional mpc synthesis, the clients' necessary requests ($?a$) would be treated as urgent. In such a scenario, the orchestration of the composition between two clients and the server (generated using the mpc synthesis algorithm) would be empty, indicating that no feasible orchestration exists.

However, if the clients' necessary requests ($?a$) are considered lazy instead, an orchestration of the composition between the server and the two clients can be achieved (computed using the orchestration synthesis). This orchestration is depicted in Figure 2. In this case, the clients take turns fulfilling their lazy necessary requests. This alternating behaviour is not possible when the necessary requests are urgent.

The orchestration in Figure 2 is obtained after three iterations of the algorithm specified in Definition 3. Initially, $\mathcal{K}_0 = \mathcal{A} = $ **Server** ⊗ **Client2** ⊗ **Client2** and $R_0 = Dangling(\mathcal{A}) = \varnothing$.

With respect to the orchestration in Figure 2, the automaton $\mathcal{A}$ contains four additional transitions that are $t_1 = [1,0,1] \xrightarrow{[-,?a,-]_\square} [1,1,1]$, $t_2 = [1,1,0] \xrightarrow{[-,-,?a]_\square} [1,1,1]$, $t_3 = [1,1,1] \xrightarrow{[!\tau,-,-]} [2,1,1]$ and $t_4 = [2,1,1] \xrightarrow{[!a,-,-]} [3,1,1]$. In the first iteration, t_1 and t_2 are removed from $\mathcal{K}_1$ because they are request transitions. We have $T_{\mathcal{K}_1} = T_{\mathcal{K}_0} \smallsetminus \{t_1, t_2\}$. Since there are no forbidden states, these are the only two transitions that are removed during the first iteration.

Concerning the set of forbidden states R_1, we have that $t_1 \in T_{\mathcal{A}}^{\square_l}$ is controllable in $\mathcal{K}_1$ via transition $[0,0,0] \xrightarrow{[a!,a?,-]_\square} [1,1,0]$. Similarly, $t_2 \in T_{\mathcal{A}}^{\square_l}$ is controllable in $\mathcal{K}_1$ via $[0,0,0] \xrightarrow{[a!,-,a?]_\square} [1,0,1]$. Hence, the source states of t_1 and t_2 will not be added to R_1. Concerning the set $Dangling(\mathcal{K}_1)$, state $[1,1,1]$ was the target of only t_1 and t_2. Moreover, state $[2,1,1]$ was the target of only t_3. Therefore, states $[1,1,1]$ and $[2,1,1]$ are unreachable in $\mathcal{K}_1$. We have that $R_1 = Dangling(\mathcal{K}_1) = \{[1,1,1],[2,1,1]\}$. In the subsequent iteration $i = 2$, since transition t_3 has target in R_1, we have $T_{\mathcal{K}_2} = T_{\mathcal{K}_1} \smallsetminus \{t_3\}$, whilst $R_2 = R_1$.

Finally, we reach the fixpoint at iteration $i = 3$, where $T_{\mathcal{K}_3} = T_{\mathcal{K}_2}$ and $R_3 = R_2$. The finalising operations for obtaining the orchestration **O** in Figure 2 from the fixpoint $\mathcal{K}_3$ consist in removing the states in R_3, i.e., $Q_{\mathbf{O}} = Q_{\mathcal{K}_3} \smallsetminus R_3$, and removing the remaining unreachable transitions in $\mathcal{K}_3$. In this case, transition $t_4 \in T_{\mathcal{K}_3}$ is removed from the orchestration, i.e., $T_{\mathbf{O}} = T_{\mathcal{K}_3} \smallsetminus \{t_4\}$.

In the subsequent section, we will delve deeper into additional details and interpretations regarding the semi-controllable transitions of contract automata.

3 Refined Semi-Controllability

We start by introducing a refined notion of semi-controllability to be used in the orchestration synthesis, formalised below. After that, we discuss how this refined notion may assist to discard some counter-intuitive orchestrations.

Definition 4 (Refined Semi-Controllability)**.** *Let $\mathcal{A}$ be an MSCA and let $t = (\vec{q}_t, \vec{a}_t, \vec{q}_t{}') \in T_{\mathcal{A}}^{\square_l}$. Moreover, given $\mathcal{A}' \subseteq \mathcal{A}$, if $\exists t' = (\vec{q}_{t'}, \vec{a}_{t'}, \vec{q}_{t'}{}') \in T_{\mathcal{A}'}^{\square_l}$ in $\mathcal{A}'$ such that the following hold:*

1. *$\vec{a}_{t'}$ is a match, $\vec{q}_{t'}, \vec{q}_{t'}{}' \notin Dangling(\mathcal{A}')$, $\vec{q}_{t(j)} = \vec{q}_{t'(j)}$, $\vec{a}_{t(j)} = \vec{a}_{t'(j)} = ?a$, for some $j \in 0 \ldots rank(\mathcal{A})$; and*

2. *there exists a sequence of transitions $t_0, \ldots, t_n$ of $\mathcal{A}'$ such that $\forall i \in 0 \ldots n$, $t_i = (\vec{q}_i, \vec{a}_i, \vec{q}_i{}')$ and the following hold:*

   - *$\vec{q}_0 = \vec{q}_t$;*
   - *$t_n = t'$;*
   - *$\vec{q}_i, \vec{q}_i{}' \notin Dangling(\mathcal{A}')$; and*
   - *if $i < n$, then $\vec{a}_{i(j)} = -$ and $\vec{q}_i{}' = \vec{q}_{i+1}$;*

then t is controllable *in $\mathcal{A}'$ (via t'). Otherwise, t is* uncontrollable *in $\mathcal{A}'$.*

By comparing Definition 2 and Definition 4, we note that only the semi-controllable transitions have been refined, whilst the others are unaltered. Conditions 1 and 2 contain the constraints that are used to decide when a semi-controllable transition is controllable or uncontrollable. The constraints of Condition 1 are also present in Definition 2. The intuition is that a (refined) semi-controllable transition t becomes controllable if (similarly to Definition 2) in a given portion of $\mathcal{A}$, there exists a semi-controllable match transition t', with source and target states not dangling, such that in both t and t' the *same* service, in the *same* local state, does the *same* request. Condition 2 of Definition 4 imposes new further constraints. It requires that t' is *reachable* from the source state of t through a sequence of transitions where the service performing the request is idle.

Consider the Venn diagram in Figure 3. The outermost set *Orchestrations* contains all orchestrations of contract automata that are computed using the notion of semi-controllability of Definition 2. The innermost set *Refined* contains only those orchestrations that are computed using the refined notion of semi-controllability in Definition 4. Intuitively, the refined notion imposes a further constraint on *when* a semi-controllable transition is controllable. As a result, more semi-controllable transitions are uncontrollable than in the previous definition. This explains why *Refined* is contained in *Orchestrations*.

All the examples of semi-controllability available in the literature [13, 8, 9, 6] (e.g., Hotel service) and Figure 2 are orchestrations belonging to the set *Refined* in Figure 3. This means that by updating the notion of semi-controllability, all orchestrations of these examples remain unaltered.
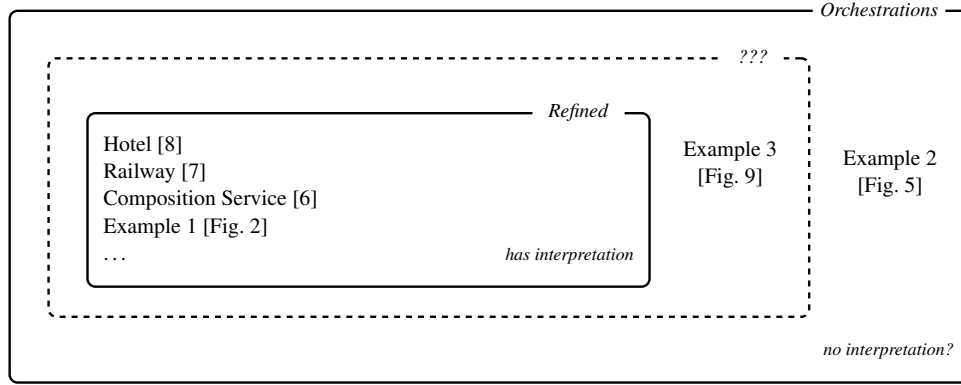
Figure 3: A Venn diagram showing the set of orchestrations of contract automata

Example 2. We now provide an example of an orchestration belonging to *Orchestrations* ∖ *Refined* (cf. Figure 3). We have three principal contracts, namely **A**lice, **B**ob and **C**arl, depicted in Figure 4. The contracts of **B**ob and **C**arl perform two alternative necessary requests. The contract of **A**lice has two branches. In each branch, a request of **B**ob and a request of **C**arl are fulfilled by corresponding offers.

Using the notion of semi-controllability from Definition 2, the synthesis algorithm of Definition 3 takes as input the composed automaton and returns the orchestration of the composition, depicted in Figure 5, which is a contract of rank 3. Indeed, for each necessary request of each service, there *exists* a match transition in the composition where the necessary request is fulfilled by a corresponding offer. In other words, for each necessary request of **B**ob and **C**arl, there exists an execution where the request is matched by a corresponding offer. For example, the composition **A**lice $\otimes$ **B**ob $\otimes$ **C**arl contains the transition $t = [a_1, b_0, c_0] \xrightarrow{[-,?d,-]_\square} [a_1, b_2, c_0]$, which is semi-controllable. According to Definition 2, t is controllable (in **A**lice $\otimes$ **B**ob $\otimes$ **C**arl) via $t' = [a_2, b_0, c_0] \xrightarrow{[!d,?d,-]_\square} [a_4, b_2, c_0]$. Since t is controllable and it is not in agreement (i.e., the label of t is a request), this transition is pruned during the synthesis of the orchestration. We note that t is controllable in t' also in all sub-automaton of the composition computed in the various iterations of the synthesis algorithm, and in the final orchestration depicted in Figure 5.

Using the refined notion of semi-controllability of Definition 4, the orchestration of **A**lice $\otimes$ **B**ob $\otimes$ **C**arl is empty (i.e., there is no orchestration). Consider again transition t. From state $[a_1, b_0, c_0]$, it is not possible to reach any transition labelled by $[!d, ?d, -]_\square$. It follows that t is uncontrollable. Hence, at some iteration i of the orchestration synthesis algorithm in Definition 3, state $[a_1, b_0, c_0]$ becomes forbidden and it is added to the set R_i. At iteration $i + 1$, the controllable transition $[a_0, b_0, c_0] \xrightarrow{[!a,-,-]_\square} [a_1, b_0, c_0]$ is pruned because its target state is forbidden. At the next iteration $(i + 2)$, the initial state $[a_0, b_0, c_0]$ becomes forbidden, because there are semi-controllable transitions not in agreement exiting the initial state (e.g., $[a_0, b_0, c_0] \xrightarrow{[-,?c,-]_\square} [a_0, b_1, c_0]$) that are uncontrollable in the sub-automaton whose transitions are T_{i+2}. Since the initial state is forbidden, it follows that there is no orchestration for **A**lice $\otimes$ **B**ob $\otimes$ **C**arl.

Indeed, whenever the state $[a_1, b_0, c_0]$ is reached, although **B**ob and **C**arl are still in their initial state, **B**ob can no longer perform the necessary request $?d$ and **C**arl can no longer perform the request $?f$. In fact, neither **B**ob nor **C**arl can decide internally which necessary request to execute from their current state. For example, there is no trace where the request $?c$ of **B**ob and the request $?f$ of **C**arl are matched.

The orchestrations belonging to *Refined* (i.e., orchestrations computed using the refined notion of semi-controllability given in Definition 4) have an intuitive interpretation when compared to the classic notion of uncontrollability. We recall that uncontrollable transitions are called *urgent* necessary transitions in MSCA, while semi-controllable transitions are called *lazy* necessary transitions.

Figure 4: Contracts of **A**lice, **B**ob and **C**arl

Figure 5: Orchestration $\mathbf{O}(\mathbf{A} \otimes \mathbf{B} \otimes \mathbf{C})$ of $\mathbf{A}$lice $\otimes \mathbf{B}$ob $\otimes \mathbf{C}$arl

Intuitively, an urgent transition cannot be delayed, whereas this is the case for a lazy one. In a concurrent composition of agents, the scheduling of concurrent urgent necessary transitions is *uncontrollable*. Instead, concerning concurrent lazy necessary transitions, each agent *internally* decides its next lazy necessary transition to execute, but the orchestrator schedules when this transition will be executed, i.e., the scheduling is *controllable*. In Example 2, there is no orchestration because, for example, from state $[a_1, b_0, c_0]$ there is no possible scheduling that allows the services to match all their necessary requests. Continuing Example 1, the orchestration in Figure 2 is non-empty because the scheduling of the actions in the orchestration is *controlled* by the orchestrator: one of the two necessary requests is scheduled to be matched only when the server has reached its internal state $[2]$. If instead the clients' necessary request $?a$ is urgent, then there exists no orchestration of the composition of two clients and the server. This is because in this case the scheduling is *uncontrollable*: it is not possible to schedule one of the two clients to have its necessary urgent request to be matched only when the server reaches the state $[2]$. In this case, the server should be ready to match the requests whenever they can be executed, without delaying them.

4 Research Challenges

In this section, we describe the currently known limits of the synthesis of orchestrations adopting either Definition 2 or Definition 4, we identify a number of research challenges to overcome these limits, and we propose a research roadmap aimed to tackle these challenges effectively.

First, the notion of semi-controllability introduced in [8, 5] and recalled in Definition 2 allows to synthesise orchestrations that may sometimes limit the capability of each service to perform internal choices. The contract automata formalism abstracts from the way that choices are made. Different implementations are possible in which each service may or may not decide the next step in an orchestration [6].

Consider again Example 2. Both **B**ob and **C**arl are able to perform two alternative necessary requests from their initial state. However, as shown in Figure 5, they are forbidden from internally deciding which necessary request is to be executed at runtime. If, for example, **B**ob selects the request $?d$ and **C**arl selects the request $?e$, then it is not possible for **A**lice to match both requests.
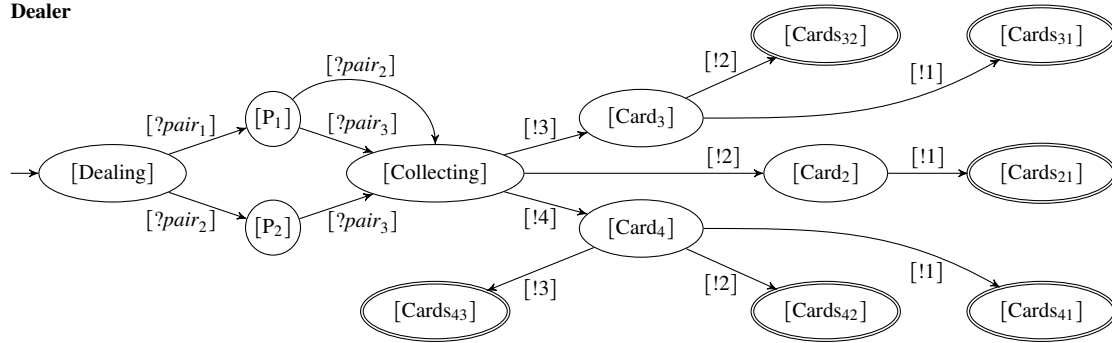
Dealer

Figure 6: Contract of the **Dealer**
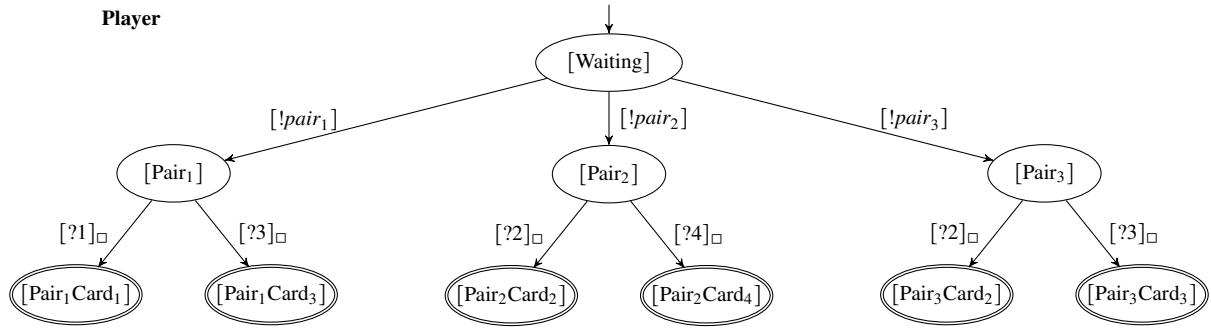
Player

Figure 7: Contract of the **Player**

If we adopt the interpretation given previously (i.e., agents internally choose their necessary transitions and their scheduling is controllable) then we argue that the orchestration computed using Definition 2 is too abstract and should in fact be empty. This is indeed the case if Definition 4 were used instead of Definition 2.

The first research challenge is to identify a concrete application of services that perform necessary requests and whose orchestration belongs to the set *Orchestrations \ Refined*.

Solving this challenge could help provide an intuitive interpretation of these types of orchestrations. An application should be identified in which each service statically requires that for each necessary request there must exist an execution where this is eventually matched (cf. Definition 2). However, during execution, the choice of which necessary request is to be matched could be external to the service performing the necessary request. Even if the execution of different branches is determined externally, a service contract may still require all branches to be available in the composition. This could be due to the contract's need to enforce certain hyperproperties, such as non-interference or opacity.

Next, we illustrate the second research challenge. All examples of orchestrations currently available in the literature [7, 6, 5, 9, 8] reside inside the set *Refined* (cf. Figure 3). We showed in Example 2 an orchestration **O** not belonging to the set *Refined* and we argued that **O** is too abstract and should in fact be empty. We now provide another example of an orchestration not belonging to the set *Refined*. However, differently from Example 2, in this case the orchestration should not be empty.

Example 3. This example involves a simple card game with two players and a dealer. At the beginning of each round, the dealer chooses a pair of cards to deal to each player (i.e., each player receives a pair of cards). The dealer can select two out of three different pairs of cards:
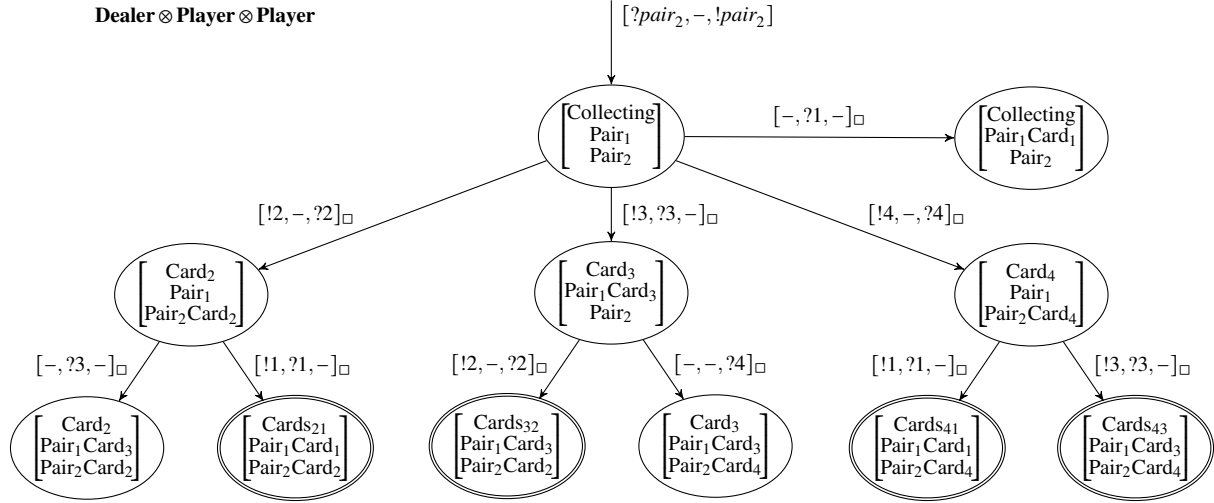
Figure 8: A fragment of the composition of **Dealer** $\otimes$ **Player** $\otimes$ **Player**

- Pair 1: card 1 and card 3;

- Pair 2: card 2 and card 4;

- Pair 3: card 2 and card 3.

After the dealer has dealt the pairs of cards, each player selects one of the two cards that was received. Once the players have selected their cards, the dealer collects the selected cards from each player. The goal of the game is for the dealer to avoid picking up two cards in ascending or equal order, which would result in the dealer losing. In other words, if the dealer picks up a card that is higher than the other card that was picked up or if two cards of the same value are picked up, the dealer loses. To ensure that the dealer never loses, the dealer has to choose the correct pairs of cards to deal. There are six possible ways to choose the pairs of cards, but only two of them guarantee a strategy for the dealer to collect the cards selected by the players in descending order. The strategy for the dealer consists of dealing to the players (in no particular order) Pair 1 and Pair 2. Indeed, in the remaining cases there exists the possibility that the players *internally* select the same card. In this case, there is no way of rearranging the transitions to avoid the same cards being picked by the dealer.

We modelled this above-mentioned problem as an orchestration of contracts, using the refined notion of semi-controllability. We only model one round of the game. The CA in Figure 6 models the dealer. Note that each request can be matched by either of the two players. Once the dealer has dealt the pairs of cards, the cards selected by the players are collected. Note that the two cards can only be collected in descending order. The CA in Figure 7 models a player. Once the player has received a card, the player decides internally which card to select. This internal decision is modelled as a choice among lazy necessary transitions.

The synthesis algorithm adopting the refined notion of semi-controllability from Definition 4 takes as input the composition of the dealer CA and two players CA and returns an empty orchestration. To explain why the resulting orchestration is empty, consider Figure 8 depicting a portion of the composition of the dealer with two players.

The state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$ is reached when the first player receives $pair_1$ and the second player receives $pair_2$. A symmetric argument holds for state $[\text{Collecting}, \text{Pair}_2, \text{Pair}_1]$, not depicted here.

The transition

$$[\text{Card}_2, \text{Pair}_1, \text{Pair}_2\text{Card}_2] \xrightarrow{[-,?3,-]_\square} [\text{Card}_2, \text{Pair}_1\text{Card}_3, \text{Pair}_2\text{Card}_2]$$

is uncontrollable according to Definition 4. Indeed, from state $[\text{Card}_2, \text{Pair}_1, \text{Pair}_2\text{Card}_2]$ it is not possible to reach state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$. This makes the state $[\text{Card}_2, \text{Pair}_1, \text{Pair}_2\text{Card}_2]$ forbidden. Hence, to avoid reaching a forbidden state, the algorithm prunes the transition

$$[\text{Collecting}, \text{Pair}_1, \text{Pair}_2] \xrightarrow{[!2,-,?2]_\square} [\text{Card}_2, \text{Pair}_1, \text{Pair}_2\text{Card}_2]$$

which is in fact controllable according to Definition 4. Indeed, from state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$ it is possible to reach the transition

$$[\text{Card}_3, \text{Pair}_1\text{Card}_3, \text{Pair}_2] \xrightarrow{[!2,-,?2]_\square} [\text{Cards}_{32}, \text{Pair}_1\text{Card}_3, \text{Pair}_2\text{Card}_2]$$

via a transition in which the second player is idle. However, during the synthesis algorithm also the state $[\text{Card}_3, \text{Pair}_1\text{Card}_3, \text{Pair}_2]$ becomes forbidden due to its outgoing necessary transition, which is uncontrollable according to Definition 2. This in turn causes the pruning of transition

$$[\text{Collecting}, \text{Pair}_1, \text{Pair}_2] \xrightarrow{[!3,?3,-]_\square} [\text{Card}_3, \text{Pair}_1\text{Card}_3, \text{Pair}_2]$$

which is controllable. Once the transition has been pruned, the transition

$$[\text{Collecting}, \text{Pair}_1, \text{Pair}_2] \xrightarrow{[!2,-,?2]_\square} [\text{Card}_2, \text{Pair}_1, \text{Pair}_2\text{Card}_2]$$

which was previously controllable becomes uncontrollable. This makes the state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$ forbidden. Note, however, that $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$ should *not* be forbidden. Indeed, from that state, for each pair of cards selected by the players, the dealer has a strategy to pick them in the correct order:

- if player 1 selects card 1 and player 2 selects card 2, then execute $[!2,-,?2],[!1,?1,-]$;

- if player 1 selects card 1 and player 2 selects card 4, then execute $[!4,-,?4],[!1,?1,-]$;

- if player 1 selects card 3 and player 2 selects card 2, then execute $[!3,?3,-],[!2,-,?2]$;

- if player 1 selects card 3 and player 2 selects card 4, then execute $[!4,-,?4],[!3,?3,-]$.

This example shows that there are cases for which Definition 4 is too restrictive. In this case, the orchestration can be computed using Definition 2, and it is displayed in Figure 9.

To better understand the underlying assumption of Definition 4, we need to decouple the moment in which a service *selects* which transition it will execute from the moment in which a service *executes* that transition. The underlying assumption of Definition 4 is that these two moments are not decoupled.

For example, the first player whose internal state is Pair_1 could select and execute ?3 also from state $[\text{Card}_2, \text{Pair}_1, \text{Pair}_2\text{Card}_2]$, while the strategy described above assumes that the player selects a card in state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$. In fact, the current implementation of the contract automata runtime environment CARE [6] allows the decoupling of these two moments. Once state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$ is reached, the orchestrator interacts with both players and, based on their choices, correctly schedules the transitions of the dealer and the players. This means that the players select their next action in state $[\text{Collecting}, \text{Pair}_1, \text{Pair}_2]$ and afterwards their execution is bounded to the transition they have selected. Summarising, Example 2 has showed that in some cases Definition 2 is too abstract, whereas Example 3 has showed that in some cases Definition 4 is too restrictive.

O(**Dealer** ⊗ **Player** ⊗ **Player**)

$\begin{bmatrix} \text{Cards}_{43} \\ \text{Pair}_1\text{Card}_3 \\ \text{Pair}_2\text{Card}_4 \end{bmatrix}$

$\begin{bmatrix} \text{Dealing} \\ \text{Waiting} \\ \text{Waiting} \end{bmatrix}$

$\begin{bmatrix} \text{Cards}_{43} \\ \text{Pair}_2\text{Card}_4 \\ \text{Pair}_1\text{Card}_3 \end{bmatrix}$

$[!3,?3,-]_\square$ $[?pair_1,!pair_1,-]$ $[?pair_1,-,!pair_1]$ $[!3,-,?3]_\square$

$\begin{bmatrix} \text{Cards}_{41} \\ \text{Pair}_1\text{Card}_1 \\ \text{Pair}_2\text{Card}_4 \end{bmatrix}$ $\begin{bmatrix} \text{Card}_4 \\ \text{Pair}_1 \\ \text{Pair}_2\text{Card}_4 \end{bmatrix}$ $\begin{bmatrix} \text{P}_1 \\ \text{Pair}_1 \\ \text{Waiting} \end{bmatrix}$ $\begin{bmatrix} \text{P}_1 \\ \text{Waiting} \\ \text{Pair}_1 \end{bmatrix}$ $\begin{bmatrix} \text{Card}_4 \\ \text{Pair}_2\text{Card}_4 \\ \text{Pair}_1 \end{bmatrix}$ $\begin{bmatrix} \text{Cards}_{41} \\ \text{Pair}_2\text{Card}_4 \\ \text{Pair}_1\text{Card}_1 \end{bmatrix}$

$[!1,?1,-]_\square$ $[!4,-,?4]_\square$ $[?pair_2,-,!pair_2]$ $[!4,?4,-]_\square$ $[!1,-,?1]_\square$

$[?pair_2,!pair_2,-]$

$\begin{bmatrix} \text{Cards}_{32} \\ \text{Pair}_1\text{Card}_3 \\ \text{Pair}_2\text{Card}_2 \end{bmatrix}$ $\begin{bmatrix} \text{Card}_3 \\ \text{Pair}_1\text{Card}_3 \\ \text{Pair}_2 \end{bmatrix}$ $\begin{bmatrix} \text{Collecting} \\ \text{Pair}_1 \\ \text{Pair}_2 \end{bmatrix}$ $\begin{bmatrix} \text{Collecting} \\ \text{Pair}_2 \\ \text{Pair}_1 \end{bmatrix}$ $\begin{bmatrix} \text{Card}_3 \\ \text{Pair}_2 \\ \text{Pair}_1\text{Card}_3 \end{bmatrix}$ $\begin{bmatrix} \text{Cards}_{32} \\ \text{Pair}_2\text{Card}_2 \\ \text{Pair}_1\text{Card}_3 \end{bmatrix}$

$[!2,-,?2]_\square$ $[!3,?3,-]_\square$ $[!3,-,?3]_\square$ $[!2,?2,-]_\square$

$[!2,-,?2]_\square$ $[!2,?2,-]_\square$

$\begin{bmatrix} \text{Cards}_{21} \\ \text{Pair}_1\text{Card}_1 \\ \text{Pair}_2\text{Card}_2 \end{bmatrix}$ $\begin{bmatrix} \text{Card}_2 \\ \text{Pair}_1 \\ \text{Pair}_2\text{Card}_2 \end{bmatrix}$ $\begin{bmatrix} \text{Card}_2 \\ \text{Pair}_2\text{Card}_2 \\ \text{Pair}_1 \end{bmatrix}$ $\begin{bmatrix} \text{Cards}_{21} \\ \text{Pair}_2\text{Card}_2 \\ \text{Pair}_1\text{Card}_1 \end{bmatrix}$

$[!1,?1,-]_\square$ $[!1,-,?1]_\square$

Figure 9: Orchestration **O**(**Dealer** ⊗ **Player** ⊗ **Player**)

> The second research challenge is to identify a notion of semi-controllability capable of discarding orchestrations such as the one in Example 2 and providing non-empty orchestrations in scenarios such as the one described in Example 3.

The resulting, currently unknown set of orchestrations that would be identified by the notion of semi-controllability that solves this challenge is depicted in Figure 3 with dashed lines.

We continue by discussing further research challenges for the orchestration synthesis of contract automata. An important aspect is the ability to scale to large orchestrations when many service contracts are composed. We note that computing Definition 4 is harder than computing Definition 2, due to the additional constraint of reachability which requires a visit of the automaton. Decoupling the moment in which a service selects a choice from the moment in which the selected choice is executed, could further increase the hardness of deciding when a lazy necessary transition is controllable.

Consider again the CA in Figure 1. From their initial state, both **B**ob and **C**arl have two choices. If, instead of two principals, we had ten principals whose behaviour is similar to that of **B**ob and **C**arl, then there would be 2^{10} possible combinations of (internal) choices the services could make.

> The third research challenge is to provide scalable solutions for synthesising orchestrations.

Generally speaking, the behaviour of an orchestration that belongs to the unknown dotted set of Figure 3 must be a sub-automaton of an orchestration computed using Definition 2 and a super-automaton of an orchestration computed using Definition 4. Indeed, Definition 2 can be used as an upper bound and Definition 4 as a lower bound to approximate the behaviour of such an orchestration.

Finally, we discuss the last research challenge identified in this paper. We previously formalised the notion of lazy necessary request that is semi-controllable according to either Definition 2 or Definition 4. We noted that Definition 2 may exclude the case in which, in the presence of a choice, a service may internally select its necessary transition. Instead, Definition 4 may exclude the case in which, in the presence of a choice, the moment in which the service internally selects its necessary transition is decoupled from the moment in which the selected necessary transition is executed. In other words, we identified two *requirements* that an orchestration of services should satisfy: *independence* and *decoupling* of choices.

> The fourth research challenge is to consolidate a set of requirements that a desirable orchestration of service contracts must satisfy.

The requirements that would solve this challenge should be established incrementally, as discussed in this paper. Formal definitions of necessary service transitions and practical examples are useful to identify the ideal set of requirements that an orchestration of services should satisfy. Of course, these requirements are entangled with the underlying execution support of an orchestration of services, which was recently proposed in [6].

4.1   Research Roadmap

We have presented a series of research challenges associated with the orchestration of contract automata. We now propose a potential research roadmap aimed at tackling these challenges effectively. However, it is necessary to further examine the concepts described below to determine their validity.

Specifying Choices   We propose to concretise the selection of the next transition to execute at contract automata level, distinguishing between internal and external selections. Presently, this distinction is abstracted away within contracts and handled by the underlying execution support. Our rationale is that abstracting from the selection process may lead to scalability challenges. Specifically, if a transition is selected internally, it must always be available, whereas an externally selected transition can be removed from the orchestration. In essence, internal selection imposes stricter requirements than external selection. Consequently, treating all selections as internal to ensure independence of choice leads to larger state spaces. For instance, the issue highlighted in Example 2 arises due to the presence of externally selected transitions. By allowing contracts to specify which transitions are internally or externally selected, we can potentially reduce the state space, as compared to considering all choices as internal.

Pursuing the above has important implications. Firstly, it necessitates updating accordingly the underlying execution support, CARE, to align it with the contract automata specifications. This entails reducing the implementation freedom for each choice to adhere to the contract's explicit selection of the next transition. By explicating choices within contracts, we establish the interpretation of necessary requests discussed in this paper. In this interpretation, a service *internally* decides to perform a necessary request, but the scheduling of the execution of the request is controlled by the orchestrator. In other words, optional actions are externally selected, whereas necessary actions are internally selected. Consequently, by explicitly stating choices in contracts, we can address the first research challenge. Indeed, all necessary requests would be internally selected. Scenarios like the one outlined in the contect of the first research challenge (i.e., external necessary requests) would be practically ruled out.

Another implication involves associating optional actions with offers and necessary actions with requests, which helps eliminate undefined choices. For instance, a branch where all actions (both offers and requests) are optional would require runtime support to determine which service is responsible for choosing the next step.

Figure 10: Two contracts whose orchestration requires further investigation

A third and final implication relates to the fourth research challenge, which entails consolidating a set of requirements for effective orchestrations. Notably, if choices are explicitly specified in contracts, the requirement of independence of choice can be removed.

Implementing the Decoupling of Choices   The second research challenge, as mentioned previously, revolves around the absence of the decoupling of choice requirement in both Definitions 2 and 4. This requirement suggests a potential implementation of semi-controllable transitions and may help identify the currently unknown set of orchestrations in Figure 3. Currently, a semi-controllable transition is defined as a transition that can be either controllable or uncontrollable based on a global condition of the automaton. However, decoupling the moment when a service internally selects a transition from the moment when the transition is executed might require splitting a semi-controllable transition into two distinct transitions.

Reasoning in this way suggests that a semi-controllable transition could potentially be represented as two consecutive transitions. The first transition would be uncontrollable, capturing the internal selection, while the subsequent transition would be controllable and responsible for executing the action. For example, consider a semi-controllable transition $[q]\xrightarrow{[?a]_{\square l}}[q']$, which would be split into two transitions: $t_1 = [q]\xrightarrow{[\tau]_{\square u}}[i]$ and $t_2 = [i]\xrightarrow{[?a]}[q']$. Here, t_1 represents an uncontrollable silent transition to an intermediate, non-final state, while t_2 is controllable and executes the action. This approach suggests that the orchestrator cannot control the internal selection made with t_1, but it can control and schedule the execution of the action indicated by t_2. Moreover, an important consequence of the fact that the intermediate state is non-final, is that t_2 must eventually be executed.

Further exploration is required to determine whether this interpretation of semi-controllability solves the third research challenge. In particular, there are still corner cases that require further investigation. For instance, consider the contracts in Figure 10. Although an orchestration could be obtained by matching the necessary request ?b of Bruce first and only afterwards the necessary request ?a of Adrian, this orchestration is not supported by the notion of semi-controllability outlined above. In this orchestration, Adrian internally selects the request ?a and the orchestrator schedules the request of Adrian to be matched later after Adrian matches the request of Bruce.

Furthermore, we envision the establishment of a clear separation between optional and necessary transitions on the one hand and controllable and uncontrollable transitions on the other. Modellers should only define contracts with requests and offers, in which all requests are considered necessary, and all offers are considered optional. Additionally, all necessary requests should be categorised as lazy/semi-controllable, thus effectively excluding urgent necessary requests from contracts. This implies that contract automata with optional and necessary transitions should be transformed into automata with solely controllable and uncontrollable transitions, which are known as plant automata in supervisory control theory. It is worth noting that all uncontrollable transitions will serve as silent moves to represent the internal selection of a necessary transition.

Experimental Validation of Performance   The third research challenge highlights the issue of scalability and proposes the adoption of Definition 2 as an upper bound for the set of orchestrations. However, it remains unclear whether the synthesis process using Definition 2 is faster compared to synthesising using the mapped plant automaton as suggested earlier. Definition 2 necessitates a visit of the automaton at each iteration of the synthesis process to determine whether a semi-controllable transition is controllable or uncontrollable. This requirement is not present in a plant automaton consisting solely of controllable and uncontrollable transitions. On the other hand, the suggested mapping approach increases the state space of the automata by introducing an additional state for each necessary transition. As a result, it is essential to conduct further experimental research to assess the effectiveness of utilising Definition 2 as an upper bound for the set of orchestrations. This research should involve measuring the performance and efficiency of the synthesis process when employing Definition 2 and comparing it with the approach based on the mapped plant automaton.

5   Conclusion

We have presented a number of research challenges related to the orchestration synthesis of contract automata. Initially, we proposed a novel refined definition of semi-controllability and compared it to the current definition through illustrative examples. We identified various sets of orchestrations, as showed in Figure 3. Additionally, we informally discussed two prerequisites that the orchestration of contracts should satisfy: independence and decoupling of choices. Furthermore, we evaluated the current formal definitions of semi-controllability based on these requirements, which generated a series of research questions regarding the orchestration synthesis of contract automata, to be addressed in future work, possibly by following the proposed research roadmap.

References

[1] Eugene Asarin, Oded Maler, Amir Pnueli & Joseph Sifakis (1998): *Controller Synthesis for Timed Automata*. *IFAC Proc. Vol.* 31(18), pp. 447–452, doi:10.1016/S1474-6670(17)42032-5.

[2] Lorenzo Bacchiani, Mario Bravetti, Marco Giunti, João Mota & António Ravara (2022): *A Java typestate checker supporting inheritance*. *Sci. Comput. Program.* 221, doi:10.1016/j.scico.2022.102844.

[3] Franco Barbanera, Ivan Lanese & Emilio Tuosto (2022): *On Composing Communicating Systems*. In Clément Aubert, Cinzia Di Giusto, Larisa Safina & Alceste Scalas, editors: *Proceedings of the 15th Interaction and Concurrency Experience (ICE'22)*, EPTCS 365, pp. 53–68, doi:10.4204/EPTCS.365.4.

[4] Davide Basile & Maurice H. ter Beek (2021): *A Clean and Efficient Implementation of Choreography Synthesis for Behavioural Contracts*. In Ferruccio Damiani & Ornela Dardha, editors: *Proceedings of the 23rd IFIP WG 6.1 International Conference on Coordination Models and Languages (COORDINATION'21)*, LNCS 12717, Springer, pp. 225–238, doi:10.1007/978-3-030-78142-2_14.

[5] Davide Basile & Maurice H. ter Beek (2022): *Contract Automata Library*. *Sci. Comput. Program.* 221, doi:10.1016/j.scico.2022.102841. Available at https://github.com/contractautomataproject/ContractAutomataLib.

[6] Davide Basile & Maurice H. ter Beek (2023): *A Runtime Environment for Contract Automata*. In Marsha Chechik, Joost-Pieter Katoen & Martin Leucker, editors: *Proceedings of the 25th International Symposium on Formal Methods (FM'23)*, *LNCS* 14000, Springer, pp. 550–567, doi:10.1007/978-3-031-27481-7_31.

[7] Davide Basile, Maurice H. ter Beek, Laura Bussi & Vincenzo Ciancia (2023): *A Toolchain for Strategy Synthesis with Spatial Properties*. *Int. J. Softw. Tools Technol. Transf.*

[8] Davide Basile, Maurice H. ter Beek, Pierpaolo Degano, Axel Legay, Gian Luigi Ferrari, Stefania Gnesi & Felicita Di Giandomenico (2020): *Controller synthesis of service contracts with variability*. *Sci. Comput. Program.* 187, doi:10.1016/j.scico.2019.102344.

[9] Davide Basile, Maurice H. ter Beek & Rosario Pugliese (2020): *Synthesis of Orchestrations and Choreographies: Bridging the Gap between Supervisory Control and Coordination of Services*. *Log. Methods Comput. Sci.* 16(2), pp. 9:1–9:29, doi:10.23638/LMCS-16(2:9)2020.

[10] Davide Basile, Pierpaolo Degano & Gian Luigi Ferrari (2016): *Automata for Specifying and Orchestrating Service Contracts*. *Log. Methods Comput. Sci.* 12(4), pp. 6:1–6:51, doi:10.2168/LMCS-12(4:6)2016.

[11] Davide Basile, Pierpaolo Degano, Gian Luigi Ferrari & Emilio Tuosto (2014): *From Orchestration to Choreography through Contract Automata*. In Ivan Lanese, Alberto Lluch Lafuente, Ana Sokolova & Hugo Torres Vieira, editors: *Proceedings of the 7th Interaction and Concurrency Experience (ICE'14)*, *EPTCS* 166, pp. 67–85, doi:10.4204/EPTCS.166.8.

[12] Davide Basile, Pierpaolo Degano, Gian Luigi Ferrari & Emilio Tuosto (2016): *Relating two automata-based models of orchestration and choreography*. *J. Log. Algebr. Methods Program.* 85(3), pp. 425–446, doi:10.1016/j.jlamp.2015.09.011.

[13] Davide Basile, Felicita Di Giandomenico, Stefania Gnesi, Pierpaolo Degano & Gian Luigi Ferrari (2017): *Specifying Variability in Service Contracts*. In: *Proceedings of the 11th International Workshop on Variability Modelling of Software-intensive Systems (VaMoS'17)*, ACM, pp. 20–27, doi:10.1145/3023956.3023965.

[14] Maurice H. ter Beek, Josep Carmona, Rolf Hennicker & Jetty Kleijn (2017): *Communication Requirements for Team Automata*. In Jean-Marie Jacquet & Mieke Massink, editors: *Proceedings of the 19th International Conference on Coordination Models and Languages (COORDINATION'17)*, *LNCS* 10319, Springer, pp. 256–277, doi:10.1007/978-3-319-59746-1_14.

[15] Maurice H. ter Beek, G. Cledou, R. Hennicker & J. Proença (2023): *Can we Communicate? Using Dynamic Logic to Verify Team Automata*. In M. Chechik, J.-P. Katoen & M. Leucker, editors: *Proceedings of the 25th International Symposium on Formal Methods (FM'23)*, *LNCS* 14000, Springer, pp. 122–141, doi:10.1007/978-3-031-27481-7_9.

[16] Athman Bouguettaya, Munindar P. Singh, Michael N. Huhns, Quan Z. Sheng, Hai Dong, Qi Yu, Azadeh Ghari Neiat, Sajib Mistry, Boualem Benatallah, Brahim Medjahed, Mourad Ouzzani, Fabio Casati, Xumin Liu, Hongbing Wang, Dimitrios Georgakopoulos, Liang Chen, Surya Nepal, Zaki Malik, Abdelkarim Erradi, Yan Wang, M. Brian Blake, Schahram Dustdar, Frank Leymann & Michael P. Papazoglou (2017): *A Service Computing Manifesto: The Next 10 Years*. *Commun. ACM* 60(4), pp. 64–72, doi:10.1145/2983528.

[17] Alberto Camacho, Meghyn Bienvenu & Sheila A. McIlraith (2019): *Towards a Unified View of AI Planning and Reactive Synthesis*. In: *Proceedings of the 29th International Conference on Automated Planning and Scheduling (ICAPS'18)*, AAAI, pp. 58–67, doi:10.1609/icaps.v29i1.3460.

[18] Christos G. Cassandras & Stéphane Lafortune (2006): *Introduction to Discrete Event Systems*. Springer, doi:10.1007/978-0-387-68612-7.

[19] CATApp. Available at `https://github.com/contractautomataproject/ContractAutomataApp`.

[20] Rüdiger Ehlers, Stéphane Lafortune, Stavros Tripakis & Moshe Y. Vardi (2017): *Supervisory control and reactive synthesis: a comparative introduction*. *Discret. Event Dyn. Syst.* 27(2), pp. 209–260, doi:10.1007/s10626-015-0223-0.

[21] Paolo Felli, Nitin Yadav & Sebastian Sardina (2017): *Supervisory Control for Behavior Composition*. *IEEE Trans. Autom. Control* 62(2), pp. 986–991, doi:10.1109/TAC.2016.2570748.

[22] Dimitrios Kouzapas, Ornela Dardha, Roly Perera & Simon J. Gay (2018): *Typechecking protocols with Mungo and StMungo: A session type toolchain for Java*. Sci. Comput. Program. 155, pp. 52–75, doi:10.1016/j.scico.2017.10.006.

[23] Michael Luttenberger, Philipp J. Meyer & Salomon Sickert (2020): *Practical synthesis of reactive systems from LTL specifications via parity games*. Acta Inform. 57(1-2), pp. 3–36, doi:10.1007/s00236-019-00349-3.

[24] Peter J. Ramadge & Walter M. Wonham (1987): *Supervisory Control of a Class of Discrete Event Processes*. SIAM J. Control Optim. 25(1), pp. 206–230, doi:10.1137/0325013.

[25] Robert E. Strom & Shaula Yemini (1986): *Typestate: A Programming Language Concept for Enhancing Software Reliability*. IEEE Trans. Softw. Eng. 12(1), pp. 157–171, doi:10.1109/TSE.1986.6312929.

[26] André Trindade, João Mota & António Ravara (2020): *Typestates to Automata and back: a tool*. In Julien Lange, Anastasia Mavridou, Larisa Safina & Alceste Scalas, editors: *Proceedings of the 13th Interaction and Concurrency Experience (ICE'20)*, EPTCS 324, pp. 25–42, doi:10.4204/EPTCS.324.4.

[27] Nobuko Yoshida, Fangyi Zhou & Francisco Ferreira (2021): *Communicating Finite State Machines and an Extensible Toolchain for Multiparty Session Types*. In Evripidis Bampis & Aris Pagourtzis, editors: *Proceedings of the 23rd International Symposium on Fundamentals of Computation Theory (FCT'21)*, LNCS 12867, Springer, pp. 18–35, doi:10.1007/978-3-030-86593-1_2.

Partially Typed Multiparty Sessions

Franco Barbanera*

Dipartimento di Matematica e Informatica, Università di Catania, Catania, Italy

`barba@dmi.unict.it`

Mariangiola Dezani-Ciancaglini

Dipartimento di Informatica, Università di Torino, Torino, Italy

`dezani@di.unito.it`

A multiparty session formalises a set of concurrent communicating participants. We propose a type system for multiparty sessions where some communications between participants can be ignored. This allows us to type some sessions with global types representing interesting protocols, which have no type in the standard type systems. Our type system enjoys Subject Reduction, Session Fidelity and "partial" Lock-freedom. The last property ensures the absence of *locks* for participants with non ignored communications. A sound and complete type inference algorithm is also discussed.

1 Introduction

The key issue in multiparty distributed systems is the composition of independent entities such that a sensible behaviour of the whole emerges from those of the components, while avoiding type errors of exchanged messages and ensuring good communication properties like Lock-freedom. MultiParty Session Types (MPST), introduced in [16, 17], are a class of choreographic formalisms for the description and analysis of such systems. Choreographic formalism are characterised by the coexistence of two distinct but related views of distributed systems: the *global* and the *local* views. The former describes the behaviour of a system as a whole, whereas the local views specify the behaviour of the single components in "isolation". Systems described by means of MPST formalisms are usually ensured (i) their overall behaviour to adhere to a given communication protocol (represented as a global type) and (ii) to enjoy particular communication properties like Lock-freedom (the specific property we focus on in the present paper).

In [3] a MPST formalism was developed for systems using synchronous communications, where global types can be assigned to multiparty sessions (parallel composition of named processes) via a type system. Typability of a multiparty session $\mathbb{M}$ by a global type G ensures that $\mathbb{M}$ behaves as described by G and is lock-free.

The property of Lock-freedom ensures that no lock is ever reached in the evolution of a system. A lock is a system's reachable configuration where a participant, which is able to perform an action, is forever prevented to do so in any possible continuation. In particular, such a configuration is called a p-lock in case the stuck participant be p. Lock-freedom – which entails Deadlock-freedom – could be however too strong to be proved in some settings, and actually useless sometimes. As a matter of fact, for particular systems, the presence of p-locks for some participants would not be problematic and would not break their specifications. Let us assume, for instance, to have a social medium where participants can

---

ask for upgrades of their communication level (i.e. the capability describing which participants they can communicate with and which sort of messages can be sent). The upgrades are granted by some particular participant u according to the particular policy of the social medium. In case u be implemented so to reply to an unbounded number of requests, it is immediate to realise that, in case all the participants get to the highest communication level, we would be in presence of a u-lock since no more level upgrade will be requested. This would not be a problem, since what we are interested in is the possibility for all the participants to progress until no communication with u is possible. From that moment on the participants other than u must be ensured to progress, but not u. This sort of circumstance is typical in clients/servers scenarios. Given a set of participants $\mathscr{P}$, we dub a system to be $\mathscr{P}$-*excluded lock-free* whenever it is p-lock free for each participant p not belonging to $\mathscr{P}$.

In this paper we present a MPST type system in the style of [3] where it is possible to derive judgments of the new shape

$$\mathsf{G} \vdash_{\mathscr{P}} \mathbb{M}$$

We say that our typing is *partial* since some communications between participants in $\mathscr{P}$ do not appear in the global type. Our type system ensures that (*a*) the communications of the participants in $\mathbb{M}$ not belonging to $\mathscr{P}$ comply with the interaction scenario represented by G and (*b*) $\mathbb{M}$ is $\mathscr{P}$-excluded lock-free.

Contributions and structure of the paper. In Section 2 we recall the calculus of multiparty sessions from [3], together with the global types. Also, we introduce the novel notion of $\mathscr{P}$-excluded Lock-freedom, that we clarify by means of an example. Section 3 is devoted to the presentation of our "partial" type system, assigning global types to multiparty sessions where some communications can be ignored. Besides, we prove the relevant properties of partially typable sessions: Subject Reduction, Session Fidelity and $\mathscr{P}$-excluded Lock-freedom. In Section 4 we discuss a sound and complete type inference algorithm for our partial type system. A section summing up our results, discussing related works and possible directions for future work concludes the paper.

2   Multiparty Sessions and Global Types

In this section we recall the calculus of multiparty sessions and the global types defined in [3]. This calculus is simpler than the original MPST calculus [16] and many of the subsequent ones. Lack of explicit channels – even if preventing the representation of session interleaving and delegation – enables us to focus on our main concerns and allows for a clear explanation of the type system we will introduce in the next section.

We use the following base sets and notation: *messages*, ranged over by $\lambda, \lambda', \ldots$; *session participants*, ranged over by p, q, r, s, u, ...; *processes*, ranged over by $P, Q, R, S, U, \ldots$; *multiparty sessions*, ranged over by $\mathbb{M}, \mathbb{M}', \ldots$; *integers*, ranged over by $i, j, l, h, k, \ldots$; *integer sets*, ranged over by $I, J, L, H, K, \ldots$.

Definition 2.1 (Processes) Processes *are defined by:*

$$P ::=^{coind} \mathbf{0} \ \mid \ \mathsf{p}!\{\lambda_i.P_i\}_{i \in I} \ \mid \ \mathsf{p}?\{\lambda_i.P_i\}_{i \in I}$$

where $I \neq \emptyset$ and $\lambda_j \neq \lambda_h$ for $j, h \in I$ and $j \neq h$.

The symbol $::=^{coind}$ in Definition 2.1 and in later definitions indicates that the productions are interpreted *coinductively*. That is, processes are possibly infinite terms. However, we assume such processes to be *regular*, i.e., with finitely many distinct sub-processes. This is done also in [7] and it allows us to adopt

in proofs the coinduction style advocated in [21] which, without any loss of formal rigour, promotes readability and conciseness.

Processes implement the communication behaviour of participants. The output process $\mathsf{p}!\{\lambda_i.P_i\}_{i\in I}$ non-deterministically chooses one message λ_i for some $i \in I$, and sends it to the participant p, thereafter continuing as P_i. Symmetrically, the input process $\mathsf{p}?\{\lambda_i.P_i\}_{i\in I}$ waits for one of the messages λ_i from the participant p, then continues as P_i after receiving it. When there is only one output we write $\mathsf{p}!\lambda.P$ and similarly for one input. We use $\mathbf{0}$ to denote the terminated process. We shall omit writing trailing $\mathbf{0}$s in processes. We denote by $\mathsf{p}\dagger\{\lambda_i.P_i\}_{i\in I}$ either $\mathsf{p}!\{\lambda_i.P_i\}_{i\in I}$ or $\mathsf{p}?\{\lambda_i.P_i\}_{i\in I}$.

In a full-fledged calculus, messages would carry values, that we avoid for the sake of simplicity; hence no selection operation over values is included in the syntax.

Definition 2.2 (Multiparty sessions) Multiparty sessions *are expressions of the shape:*

$$\mathsf{p}_1[P_1] \parallel \cdots \parallel \mathsf{p}_n[P_n]$$

where $\mathsf{p}_j \neq \mathsf{p}_h$ *for* $1 \leq j, h \leq n$ *and* $j \neq h$. *We use* $\mathbb{M}$ *to range over multiparty sessions.*

Multiparty sessions (sessions, for short) are parallel compositions of located processes of the form $\mathsf{p}[P]$, each enclosed within a different participant p. We assume the standard structural congruence $\equiv$ on multiparty sessions, stating that parallel composition is associative and commutative and has neutral elements $\mathsf{p}[\mathbf{0}]$ for any p. If $P \neq \mathbf{0}$ we write $\mathsf{p}[P] \in \mathbb{M}$ as short for $\mathbb{M} \equiv \mathsf{p}[P] \parallel \mathbb{M}'$ for some $\mathbb{M}'$. This abbreviation is justified by the associativity and commutativity of parallel composition.

The *set of active participants* (participants for short) of a session $\mathbb{M}$, notation $\mathsf{prt}(\mathbb{M})$, is as expected:

$$\mathsf{prt}(\mathbb{M}) = \{\mathsf{p} \mid \mathsf{p}[P] \in \mathbb{M}\}$$

It is easy to verify that the sets of participants of structurally congruent sessions coincide.

To define the *synchronous operational semantics* of sessions we use an LTS, whose transitions are decorated by labels denoting message exchanges.

Definition 2.3 (LTS for Multiparty Sessions) *The* labelled transition system (LTS) for multiparty sessions *is the closure under structural congruence of the reduction specified by the unique rule:*

$$[\text{Comm-T}] \; \frac{h \in I \subseteq J}{\mathsf{p}[\mathsf{q}!\{\lambda_i.P_i\}_{i\in I}] \parallel \mathsf{q}[\mathsf{p}?\{\lambda_j.Q_j\}_{j\in J}] \parallel \mathbb{M} \xrightarrow{\mathsf{p}\lambda_h\mathsf{q}} \mathsf{p}[P_h] \parallel \mathsf{q}[Q_h] \parallel \mathbb{M}}$$

Rule [Comm-T] makes communications possible, by describing when a participant p can send a message λ_h to participant q, and what is the effect of such message exchange. This rule is non-deterministic in the choice of messages. The condition $I \subseteq J$ ensures that the sender can freely choose the message, since the receiver must offer all sender messages and possibly more. This allows us to distinguish in the operational semantics between internal and external choices. Note that this condition will be ensured by the typing Rule [Comm] (see Definition 3.1).

Let Λ range over *labels*, namely triples of the form $\mathsf{p}\lambda\mathsf{q}$. We define *traces* as (possibly infinite) sequences of labels by:

$$\sigma ::=^{coind} \varepsilon \mid \Lambda \cdot \sigma$$

where ε is the empty sequence. We use $|\sigma|$ to denote the length of the trace σ, where $|\sigma| = \infty$ when σ is an infinite trace. We define the participants of labels and traces:

$$\mathsf{prt}(\mathsf{p}\lambda\mathsf{q}) = \{\mathsf{p},\mathsf{q}\} \qquad \mathsf{prt}(\varepsilon) = \emptyset \qquad \mathsf{prt}(\Lambda \cdot \sigma) = \mathsf{prt}(\Lambda) \cup \mathsf{prt}(\sigma)$$

When $\sigma = \Lambda_1 \cdot \ldots \cdot \Lambda_n$ ($n \geq 0$) we write $\mathbb{M} \xrightarrow{\sigma} \mathbb{M}'$ as short for $\mathbb{M} \xrightarrow{\Lambda_1} \mathbb{M}_1 \cdots \xrightarrow{\Lambda_n} \mathbb{M}_n = \mathbb{M}'$. As usual we write $\mathbb{M} \rightarrow$ (resp. $\mathbb{M} \nrightarrow$) when there exist (resp. no) Λ and $\mathbb{M}'$ such that $\mathbb{M} \xrightarrow{\Lambda} \mathbb{M}'$.

It is easy to verify that, in a transition, only the two participants of its label are involved, as formalised below.

Fact 2.4 *If* $\{p, q\} \cap \{r, s\} = \emptyset$ *and* $r[R] \parallel s[S] \parallel \mathbb{M} \xrightarrow{p\lambda q} r[R] \parallel s[S] \parallel \mathbb{M}'$, *then*

$$r[R'] \parallel s[S'] \parallel \mathbb{M} \xrightarrow{p\lambda q} r[R'] \parallel s[S'] \parallel \mathbb{M}'$$

for arbitrary R', S'.

We define now the property of $\mathscr{P}$-excluded Lock-freedom, a "partial" version of the standard Lock-freedom [19, 26]. The latter consists in the possible eventual completion of pending communications of any participant (this can be alternatively stated by saying that any participant is lock-free). We are interested instead in the progress of some specific participants only, namely those we decide not to "ignore". In the following, $\mathscr{P}$ will range over sets of ignored participants.

Definition 2.5 ($\mathscr{P}$-excluded Lock-freedom) *A multiparty session* $\mathbb{M}$ *is a* $\mathscr{P}$-excluded lock-free *session if* $\mathbb{M} \xrightarrow{\sigma} \mathbb{M}'$ *and* $p \in \text{prt}(\mathbb{M}') \setminus \mathscr{P}$ *imply* $\mathbb{M}' \xrightarrow{\sigma' \cdot \Lambda} \mathbb{M}''$ *for some* σ' *and* Λ *such that* $p \in \text{prt}(\Lambda)$.

It is natural to extend also the usual notion of Deadlock-freedom to our setting.

Definition 2.6 ($\mathscr{P}$-excluded Deadlock-freedom) *A multiparty session* $\mathbb{M}$ *is a* $\mathscr{P}$-excluded deadlock-free *session if* $\mathbb{M} \xrightarrow{\sigma} \mathbb{M}' \nrightarrow$ *implies* $\text{prt}(\mathbb{M}') \subseteq \mathscr{P}$.

It is immediate to check that, as for standard Lock- and Deadlock-freedom, the following hold.

Fact 2.7 *$\mathscr{P}$-excluded Lock-freedom implies $\mathscr{P}$-excluded Deadlock-freedom.*

The vice versa does not hold. For example if $P = q!\lambda.P$, $Q = p?\lambda.P$ and $R \neq \mathbf{0}$, then $p[P] \parallel q[Q] \parallel r[R]$ is $\mathscr{P}$-excluded deadlock-free for any $\mathscr{P}$, but $\mathscr{P}$-excluded lock-free only when $r \in \mathscr{P}$.

The following example illustrates the notion of $\mathscr{P}$-excluded Lock-freedom.

Example 2.8 (Social media) Let us consider a system describing a simplified social media situation. Participant q is allowed to greet participant p by sending a message HELLO. Participant p would like to reply to q, but in order to do that she needs to be granted a higher communication level. The task of granting permissions is performed by participant u which, upon p's request (REQ), decides – according to some parameters – whether the permission is granted (GRTD) or denied (DND). Her decision is communicated to both p and q. We assume that (for reusability motivation) u is implemented in order to process an unbounded number of requests. For what concerns p, however, once she is granted the higher communication level, she can return the greeting to q, so ending their interaction. The above system corresponds to the following session

$$\mathbb{M} \equiv p[P] \parallel q[Q] \parallel u[U]$$

where $P = q?\text{HELLO}.u!\text{REQ}.u?\{\text{DND}.P, \text{GRTD}.q!\text{HELLO}\}$, $Q = p!\text{HELLO}.u?\{\text{DND}.Q, \text{GRTD}.p?\text{HELLO}\}$ and $U = p?\text{REQ}.p!\{\text{DND}.q!\text{DND}.U, \text{GRTD}.q!\text{GRTD}.U\}$. The session is $\{u\}$-excluded lock-free, since, once p has been granted the higher communication level, we get

$$p[\mathbf{0}] \parallel q[\mathbf{0}] \parallel u[U]$$

where u is willing to interact but she will never be able. This, however, should not be deemed a problem, since we are actually interested in that the interactions between p and q do proceed smoothly.   ◇

The behaviour of multiparty sessions can be disciplined by means of types. Global types describe the conversation scenarios of multiparty sessions, possibly in a partial way.

Definition 2.9 (Global types) Global types *are defined by:*

$$G ::=^{coind} \text{End} \mid p \to q : \{\lambda_i.G_i\}_{i \in I}$$

where $I \neq \emptyset$ and $\lambda_j \neq \lambda_h$ for $j,h \in I$ and $j \neq h$.

As for processes, we allow only *regular* global types. The type $p \to q : \{\lambda_i.G_i\}_{i \in I}$ formalises a protocol where participant p must send to q a message λ_j for some $j \in I$ (and q must receive it) and then, depending on which λ_j was chosen by p, the protocol continues as G_j. We write $p \to q : \lambda.G$ when there is only one message. We use End to denote the terminated protocol. We shall omit writing trailing Ends in global types.

   We define the *set of paths of a global type*, notation *paths*(G), as the greatest set of traces such that:

$$paths(\text{End}) = \{\varepsilon\} \qquad paths(p \to q : \{\lambda_i.G_i\}_{i \in I}) = \bigcup_{i \in I}\{p\lambda_i q \cdot \sigma \mid \sigma \in paths(G_i)\}$$

The *set of participants of a global type* is the set of participants of its paths:

$$\text{prt}(G) = \bigcup_{\sigma \in paths(G)} \text{prt}(\sigma)$$

The regularity of global types ensures that such sets of participants are always finite.

   Boundedness is a property of global types that will enable us to get $\mathscr{P}$-excluded Lock-freedom from typability. This consists in requiring any participant of a global type to occur either in all the paths or in no path of any of its subterms which are global types. Notably this condition is a form of fairness, even if it strongly differs from the notions of fairness discussed in [15], where fairness assumptions rule out computational paths. Technically, we shall use the notions of *depth* and of *boundedness* as defined below. We denote by $\sigma[n]$ with $n \in \mathbb{N}$ the *n*-th label in the path σ, where $1 \leq n \leq |\sigma|$.

Definition 2.10 (Depth) *Let G be a global type. For $\sigma \in paths(G)$ we define*

$$depth(\sigma,p) = \text{infimum}\{n \mid p \in \text{prt}(\sigma[n])\}$$

and define depth(G,p), the depth *of p in G, as follows:*

$$depth(G,p) = \begin{cases} \text{supremum}\{depth(\sigma,p) \mid \sigma \in paths(G)\} & \text{if } p \in \text{prt}(G) \\ 0 & \text{otherwise} \end{cases}$$

   Note that $depth(G,p) = 0$ iff $p \notin \text{prt}(G)$. Moreover, if $p \in \text{prt}(G)$, but for some $\sigma \in paths(G)$ it is the case that $p \notin \text{prt}(\sigma[n])$ for all $n \leq |\sigma|$, then $depth(\sigma,p) = \text{infimum} \emptyset = \infty$. Hence, if p is a participant of a global type G and there is some path in G where p does not occur, then $depth(G,p) = \infty$.

Definition 2.11 (Boundedness) *A global type G is* bounded *if $depth(G',p)$ is finite for all participants $p \in \text{prt}(G')$ and all types G' which occur in G.*

   Intuitively, this means that if $p \in \text{prt}(G')$ for a type G' which occurs in G, then the search for an interaction of the shape $p\lambda q$ or $q\lambda p$ along a path $\sigma \in paths(G')$ terminates (and recall that G' can be infinite, in which case G is such). Hence the name.

   Example 2 of [3] shows the necessity of considering all types occurring in a global type when defining boundedness and that also a finite global type can be unbounded.

   Since global types are regular, the boundedness condition is decidable. We shall allow only bounded global types in typing sessions.

Example 2.12 (A global type for the social media example) The intended overall behaviour of the multiparty session $\mathbb{M}$ in Example 2.8, up to the point where the request of p is possibly accepted by u, is described by the following global type G.

$$G = q \to p{:}\text{HELLO}. \, p \to u{:}\text{REQ}. \, u \to p{:} \begin{cases} \text{DND}. \, u \to q{:}\text{DND}. \, G \\ \text{GRTD}. \, u \to q{:}\text{GRTD}. \, p \to q{:}\text{HELLO} \end{cases}$$

Typability of $\mathbb{M}$ with G – by means of the type system defined in the next section – will ensure (see Theorems 3.6 and 3.7 below) that the behaviours of participants of G, but u, will perfectly adhere to what G describes. $\diamond$

We conclude this section by defining the standard LTS for global types. By means of such LTS we formalise the intended meaning of global types as overall (possibly partial) descriptions of sessions' behaviours. It will be used in the next section to prove the properties of Subject Reduction and Session Fidelity which, in our setting, will slightly differ from the standard ones [16, 17].

Definition 2.13 (LTS for Global Types) *The* labelled transition system (LTS) for global types *is specified by the following axiom and rule:*

$$[\textsc{Ecomm}] \; \frac{}{p \to q : \{\lambda_i.G_i\}_{i \in I} \xrightarrow{\;p\lambda_j q\;} G_j} \quad j \in I$$

$$[\textsc{Icomm}] \; \frac{G_i \xrightarrow{\;p\lambda q\;} G_i' \quad \forall i \in I \qquad \{p,q\} \cap \{r,s\} = \emptyset}{r \to s : \{\lambda_i.G_i\}_{i \in I} \xrightarrow{\;p\lambda q\;} r \to s : \{\lambda_i.G_i'\}_{i \in I}}$$

Axiom [ECOMM] formalises the fact that, in a session exposing the behaviour $p \to q : \{\lambda_i.G_i\}_{i \in I}$, there are participants p and q ready to exchange a message λ_j for any $j \in I$, the former as sender and the latter as receiver. If such a communication is actually performed, the resulting session will expose the behaviour G_j.

Rule [ICOMM] makes sense since, in a global type $r \to s : \{\lambda_i.G_i\}_{i \in I}$, communications involving participants p and q, ready to interact with each other uniformly in all branches, can be performed if neither of them is involved in a previous interaction between r and s. In this case, the interaction between p and q is independent of the choice of r, and may be executed before it.

3   Type System and its Properties

As in [3, 9, 13], our type assignment allows for a simple treatment of many technical issues, by avoiding projections, local types and subtyping [16, 17]. The novelty of the type system we present in this section with respect to those in [3, 9, 13] is that the judgments are parametrised by a set $\mathscr{P}$ of participants. These are the participants whose Lock-freedom we do not care about. The simplicity of our calculus allows us to formulate a type system deriving directly global types for multiparty sessions, i.e. judgments of the form $G \vdash_{\mathscr{P}} \mathbb{M}$ (where G is bounded). Here and in the following the double line indicates that the rules are interpreted coinductively [27, Chapter 21].

Definition 3.1 (Type system) *The type system $\vdash_{\mathscr{P}}$ is defined by the following axiom and rules, where sessions are considered modulo structural congruence:*

$$[\text{End}] \; \text{End} \vdash_{\emptyset} p[\mathbf{0}]$$

$$[\textsc{Comm}] \; \frac{\begin{array}{c} G_i \vdash_{\mathscr{P}_i} p[P_i] \parallel q[Q_i] \parallel \mathbb{M} \\ (\text{prt}(G_i) \cup \mathscr{P}_i) \setminus \{p,q\} = \text{prt}(\mathbb{M}) \quad \forall i \in I \end{array}}{G \vdash_{\mathscr{P}} p[q!\{\lambda_i.P_i\}_{i \in I}] \parallel q[p?\{\lambda_j.Q_j\}_{j \in J}] \parallel \mathbb{M}} \quad \begin{array}{l} G = p \to q : \{\lambda_i.G_i\}_{i \in I} \quad G \text{ is bounded} \\ \mathscr{P} = \bigcup_{i \in I}\mathscr{P}_i \quad I \subseteq J \end{array}$$

$$[\text{WEAK}] \ \frac{\mathsf{G} \vdash_{\mathscr{P}_1} \mathbb{M}_1}{\mathsf{G} \vdash_{\mathscr{P}_1 \cup \mathscr{P}_2} \mathbb{M}_1 \parallel \mathbb{M}_2} \quad \mathscr{P}_2 = \mathsf{prt}(\mathbb{M}_2) \neq \emptyset$$

Axiom [End] simply states that the null session has the behaviour described by End. In the null session there is obviously no participant whose Lock-freedom we do not care about, hence the $\emptyset$ subscript. An alternative and sound version of this axiom could be $\mathtt{End} \vdash_{\mathsf{prt}(\mathbb{M})} \mathbb{M}$ for any $\mathbb{M}$. Such a judgment however can be easily derived using Rule [WEAK].

Rule [COMM] just adds simultaneous communications to global types and to corresponding processes inside sessions. Since the set $\mathscr{P}_i$ contains the ignored participants in branch i, the possibly non lock-free participants in the conclusion must be $\bigcup_{i \in I} \mathscr{P}_i$. Note that Rule [COMM] allows more inputs than corresponding outputs, in agreement with the condition in Rule [COMM-T]. It also allows more branches in the input process than in the global type, just mimicking the subtyping for session types [14]. Instead, the number of branches in the output process and the global type must be the same. This does not restrict typability, while it improves Session Fidelity (by arguing, respectively, as in [4] and [3]). The condition $(\mathsf{prt}(\mathsf{G}_i) \cup \mathscr{P}_i) \setminus \{\mathsf{p}, \mathsf{q}\} = \mathsf{prt}(\mathbb{M})$ for all $i \in I$ ensures that the participants in the session are exactly those we keep track of either in G and/or in $\mathscr{P}$. This condition prevents, for example, to derive $\mathsf{G} \vdash_{\emptyset} \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{r}[R]$, where $\mathsf{G} = \mathsf{p} \rightarrow \mathsf{q} : \lambda.\mathsf{G}$, $P = \mathsf{q}!\lambda.P$, $Q = \mathsf{p}?\lambda.Q$ and $R \neq \mathbf{0}$ is arbitrary. Note that, instead, it is possible to derive $\mathsf{G} \vdash_{\{\mathsf{r}\}} \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{r}[R]$ with $R \neq \mathbf{0}$ arbitrary. Lock-freedom can be ensured only for the participants of $\mathbb{M}$ not belonging to $\mathscr{P}$.

Rule [WEAK] enables to type check just a sub session as far as we do not care about the Lock-freedom of the participants of the rest of the session. We keep track of such participants in the subscript of the entailment symbol. The condition $\mathscr{P}_2 \neq \emptyset$ forbids infinite applications of this rule. This condition allows us to use coinduction on typing derivations.

Sessions are considered modulo structural congruence in typing rules following [3, 9, 13]. Clearly this could be avoided by adding an obvious typing rule, but we prefer to have a lighter type system.

The regularity of processes and global types ensures the decidability of type checking.

Example 3.2 (Typing of the social media) Let P, Q and U be defined as in Example 2.8 and

$$P_1 = \mathsf{u}!\text{REQ}.\mathsf{u}?\{\text{DND}.P, \text{GRTD}.P_2\}, P_2 = \mathsf{q}!\text{HELLO}, Q_1 = \mathsf{u}?\{\text{DND}.Q, \text{GRTD}.Q_2\}, Q_2 = \mathsf{p}?\text{HELLO}$$
$$U_1 = \mathsf{p}!\{\text{DND}.\mathsf{q}!\text{DND}.U, \text{GRTD}.\mathsf{q}!\text{GRTD}.U\}$$

Moreover, let G be defined as in Example 2.12 and $\mathsf{G}_1 = \mathsf{p} \rightarrow \mathsf{u}:\text{REQ}. \mathsf{G}_2$

$$\mathscr{D} = \cfrac{\cfrac{\mathscr{D}}{\mathsf{G}_3 \vdash_{\{\mathsf{u}\}} \mathsf{p}[P] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[\mathsf{q}!\text{DND}.U]} \qquad \cfrac{\cfrac{\cfrac{\cfrac{}{\mathtt{End} \vdash_{\emptyset} \mathsf{p}[\mathbf{0}] \parallel \mathsf{q}[\mathbf{0}]} \ [\text{End}]}{\mathsf{G}_5 \vdash_{\emptyset} \mathsf{p}[P_2] \parallel \mathsf{q}[Q_2]}}{\mathsf{G}_5 \vdash_{\{\mathsf{u}\}} \mathsf{p}[P_2] \parallel \mathsf{q}[Q_2] \parallel \mathsf{u}[U]} \ [\text{WEAK}]}{\mathsf{G}_4 \vdash_{\{\mathsf{u}\}} \mathsf{p}[P_2] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[\mathsf{q}!\text{GRTD}.U]}}{\cfrac{\cfrac{\mathsf{G}_2 \vdash_{\{\mathsf{u}\}} \mathsf{p}[\mathsf{u}?\{\text{DND}.P, \text{GRTD}.P_2\}] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_1]}{\mathsf{G}_1 \vdash_{\{\mathsf{u}\}} \mathsf{p}[P_1] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U]}}{\mathsf{G} \vdash_{\{\mathsf{u}\}} \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{u}[U]}}$$

Figure 1: A type derivation for the social media.

$$G_2 = u \to p: \begin{cases} \text{DND.} \, G_3 \\ \text{GRTD.} \, G_4 \end{cases} \qquad G_3 = u \to q:\text{DND.} \, G \qquad G_4 = u \to q:\text{GRTD.} \, G_5 \qquad G_5 = p \to q:\text{HELLO}$$

Figure 1 shows a derivation of the global type G of Example 2.12 for the multiparty session $\mathbb{M}$ of Example 2.8. The missing rule names are all [COMM]. Note how in the leftmost branch of the derivation it is possible to get $\{u\}$ as subscript without recurring to Rule [WEAK] thanks to the infiniteness of the branch. For the same motivation, in case we had $P = q?\text{HELLO.} u!\text{REQ.} u?\{\text{DND.} P, \text{GRTD.} P'\}$ with $P' = q!\text{HELLO.} P'$ and $Q = p!\text{HELLO.} u?\{\text{DND.} Q, \text{GRTD.} Q'\}$ with $Q' = p?\text{HELLO.} Q'$ (namely in case p and q kept on indefinitely exchanging HELLO messages after receiving the GRTD message) the whole resulting session would be typable without recurring to Rule [WEAK]. $\diamond$

We note that session participants are of three different kinds in a typing judgment:

1. the lock-free participants which behave as pointed out by the global type; these participants occur in the global type but do not belong to the set of ignored participants;

2. the participants which "partially" behave as pointed out by the global type and can get stuck; these participants occur in the global type and belong to the set of ignored participants;

3. the participants which behave in an unpredictable way; these participants do not occur in the global type but belong to the set of ignored participants.

We observe also that $\vdash_\emptyset$ coincides with the typing relation of [3].

In the remainder of this section we will show the main properties of our type system, i.e. Subject Reduction, Session Fidelity and $\mathscr{P}$-excluded Lock-freedom. We start with some lemmas which are handy for the subsequent proofs. All proofs are by coinduction on $G \vdash_\mathscr{P} \mathbb{M}$ and by cases on the last applied rule.

The first lemma states that, when $G \vdash_\mathscr{P} \mathbb{M}$, all participants of $\mathbb{M}$ must be participants of G and/or must belong to the set $\mathscr{P}$.

Lemma 3.3 $G \vdash_\mathscr{P} \mathbb{M}$ *implies* $\text{prt}(G) \cup \mathscr{P} = \text{prt}(\mathbb{M})$.

Proof. Rule [COMM]. Immediate by the condition $(\text{prt}(G_i) \cup \mathscr{P}_i) \setminus \{p, q\} = \text{prt}(\mathbb{M})$ for all $i \in I$.

Rule [WEAK]. In such a case, $\mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $\mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2$. By coinduction we get $\text{prt}(G) \cup \mathscr{P}_1 = \text{prt}(\mathbb{M}_1)$. We hence get the thesis by the condition $\mathscr{P}_2 = \text{prt}(\mathbb{M}_2)$. $\square$

By the above lemma, from $G \vdash_\mathscr{P} \mathbb{M}$ it is immediate to get also that $p \in \text{prt}(G)$ implies $p \in \text{prt}(\mathbb{M})$ and that $p \in \text{prt}(\mathbb{M})$ and $p \notin \mathscr{P}$ imply $p \in \text{prt}(G)$.

The process of a participant which does not occur in the global type can be freely replaced, since typing ensures nothing about the behaviour of this participant.

Lemma 3.4 *If* $G \vdash_\mathscr{P} p[P] \parallel \mathbb{M}$ *and* $p \in \text{prt}(\mathbb{M}) \setminus \text{prt}(G)$, *then* $G \vdash_{\mathscr{P}'} p[P'] \parallel \mathbb{M}$ *with* $\mathscr{P}' \subseteq \mathscr{P}$ *for an arbitrary* P'.

Proof. Rule [COMM]. Then $G = q \to r : \{\lambda_i.G_i\}_{i \in I}$ and $\mathbb{M} \equiv q[r!\{\lambda_i.Q_i\}_{i \in I}] \parallel r[q?\{\lambda_j.R_j\}_{j \in J}] \parallel \mathbb{M}_0$ and $I \subseteq J$ and $G_i \vdash_{\mathscr{P}_i} p[P] \parallel q[Q_i] \parallel r[R_i] \parallel \mathbb{M}_0$ for all $i \in I$ with $\mathscr{P} = \bigcup_{i \in I} \mathscr{P}_i$. By coinduction we get $G_i \vdash_{\mathscr{P}'_i} p[P'] \parallel q[Q_i] \parallel r[R_i] \parallel \mathbb{M}_0$ with $\mathscr{P}'_i \subseteq \mathscr{P}_i$ for all $i \in I$ for an arbitrary P'. We conclude using Rule [COMM].

Rule [WEAK]. Then $\mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2$ and $p[P] \parallel \mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $G \vdash_{\mathscr{P}_1} \mathbb{M}_1$. If $\mathbb{M}_1 \equiv p[P] \parallel \mathbb{M}'_1$ by coinduction $G \vdash_{\mathscr{P}'_1} p[P'] \parallel \mathbb{M}'_1$ with $\mathscr{P}'_1 \subseteq \mathscr{P}_1$ for arbitrary P' and we conclude using Rule [WEAK]. If $\mathbb{M}_2 \equiv p[P] \parallel \mathbb{M}'_2$ we can apply Rule [WEAK] to $G \vdash_{\mathscr{P}_1} \mathbb{M}_1$ and $p[P'] \parallel \mathbb{M}'_2$ for arbitrary P'. $\square$

Note that in previous lemma $\mathscr{P}' = \mathscr{P}$ unless $P' = \mathbf{0}$ and in this case $\mathscr{P}' \cup \{p\} = \mathscr{P}$.

If $p[q \dagger \{\lambda_i.P_i\}_{i \in I}] \in \mathbb{M}$ we say that q is the *top partner* of p and we write $tp(\mathbb{M}, p) = q$. Note that we can have $q \notin prt(\mathbb{M})$ or $tp(\mathbb{M}, q) \neq p$. For example, if $\mathbb{M} \equiv p[q \dagger \{\lambda_i.P_i\}_{i \in I}] \parallel q[r \dagger \{\lambda'_j.Q_j\}_{j \in J}]$, we have $tp(\mathbb{M}, p) = q$ and $tp(\mathbb{M}, q) = r \neq p$.

Typing ensures that if a participant occurs in a global type then also her top partner occurs in the global type.

Lemma 3.5 *If* $G \vdash_{\mathscr{P}} \mathbb{M}$ *and* $p \in prt(G)$*, then* $tp(\mathbb{M}, p) \in prt(G)$*.*

Proof. By Lemma 3.3 and $p \in prt(G)$ we have that $p \in prt(\mathbb{M})$ and then $tp(\mathbb{M}, p)$ is defined. So, let $tp(\mathbb{M}, p) = q$.

Rule [COMM]. Then $G = r \to s : \{\lambda_i.G_i\}_{i \in I}$ and $\mathbb{M} \equiv r[s!\{\lambda_i.R_i\}_{i \in I}] \parallel s[r?\{\lambda_j.S_j\}_{j \in J}] \parallel \mathbb{M}_0$ with $I \subseteq J$ and $G_i \vdash_{\mathscr{P}_i} r[R_i] \parallel s[S_i] \parallel \mathbb{M}_0$ for all $i \in I$ with $\mathscr{P} = \bigcup_{i \in I} \mathscr{P}_i$. If $p \in \{r,s\}$, then $\{p,q\} = \{r,s\}$ and we are done. Otherwise $tp(\mathbb{M}, p) = q$ implies $tp(r[R_i] \parallel s[S_i] \parallel \mathbb{M}_0, p) = q$ for all $i \in I$. Moreover $p \in prt(G)$ implies $p \in prt(G_i)$ for all $i \in I$ since G is bounded. By coinduction we get $q \in prt(G_i)$ for all $i \in I$. We conclude $q \in prt(G)$.

Rule [WEAK]. Then $\mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2$, $\mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $G \vdash_{\mathscr{P}_1} \mathbb{M}_1$. Since by Lemma 3.3 $p \in prt(G)$ implies $p \in prt(\mathbb{M}_1)$ we have $tp(\mathbb{M}_1, p) = q$. We get by coinduction $q \in prt(G)$. $\square$

In our particular setting, what Subject Reduction ensures depends on which participants we consider (unlike its standard version, e.g. in [4] and [3]). In particular, it ensures that, when the involved participants occur in the global types, the transitions of well-typed sessions are mimicked by those of global types (namely they proceed as prescribed by the global type). Otherwise the reduced session can be typed by the same global type. Key for this proof is Lemma 3.5, which ensures that the communicating participants either both occur or both do not occur in the global type.

Theorem 3.6 (Subject Reduction) *Let* $G \vdash_{\mathscr{P}} \mathbb{M}$ *and* $\mathbb{M} \xrightarrow{p\lambda q} \mathbb{M}'$*.*

    i) *If* $\{p,q\} \subseteq prt(G)$*, then* $G \xrightarrow{p\lambda q} G'$ *and* $G' \vdash_{\mathscr{P}'} \mathbb{M}'$ *with* $\mathscr{P}' \subseteq \mathscr{P}$*.*

    ii) *If* $p, q \notin prt(G)$*, then* $G \vdash_{\mathscr{P}'} \mathbb{M}'$ *with* $\mathscr{P}' \subseteq \mathscr{P}$*.*

Proof. From $\mathbb{M} \xrightarrow{p\lambda q} \mathbb{M}'$ we get $\mathbb{M} \equiv p[q!\{\lambda_i.P_i\}_{i \in I}] \parallel q[p?\{\lambda_j.Q_j\}_{j \in J}] \parallel \mathbb{M}_0$ with $I \subseteq J$ and $\mathbb{M}' \equiv p[P_l] \parallel q[Q_l] \parallel \mathbb{M}_0$ and $\lambda = \lambda_l$ for some $l \in I$. Note that Lemma 3.5 implies either $\{p,q\} \subseteq prt(G)$ or $p, q \notin prt(G)$.

ii). In this case $G \vdash_{\mathscr{P}} \mathbb{M}$ implies $G \vdash_{\mathscr{P}'} p[P_l] \parallel q[Q_l] \parallel \mathbb{M}_0$ with $\mathscr{P}' \subseteq \mathscr{P}$ by Lemma 3.4.

i). The proof is by coinduction on $G \vdash_{\mathscr{P}} p[P] \parallel \mathbb{M}$ and by cases on the last applied rule.

Rule [COMM]. We get $G = r \to s : \{\lambda'_h.G_h\}_{h \in H}$ and $\mathbb{M} \equiv r[s!\{\lambda'_h.R_h\}_{h \in H}] \parallel s[r?\{\lambda'_k.S_k\}_{k \in K}] \parallel \mathbb{M}_1$ and $H \subseteq K$ and $G_h \vdash_{\mathscr{P}_h} r[R_h] \parallel s[S_h] \parallel \mathbb{M}_1$ for all $h \in H$ with $\mathscr{P} = \bigcup_{h \in H} \mathscr{P}_h$. If with $p = r$ and $q = s$, then $I = H$, $J = K$ and $\lambda_i = \lambda'_i$ for all $i \in I$. We conclude $G \xrightarrow{p\lambda q} G_l$ and $G_l \vdash_{\mathscr{P}_l} \mathbb{M}'$. Otherwise $\{p,q\} \cap \{r,s\} = \emptyset$, which implies $r[R_h] \parallel s[S_h] \parallel \mathbb{M}_1 \xrightarrow{p\lambda q} r[R_h] \parallel s[S_h] \parallel \mathbb{M}'_1$ for all $h \in H$ by Fact 2.4. Moreover $\{p,q\} \subseteq prt(G_h)$ for all $h \in H$ since G is bounded. By coinduction we get $G_h \xrightarrow{p\lambda q} G'_h$ and $G'_h \vdash_{\mathscr{P}'_h} r[R_h] \parallel s[S_h] \parallel \mathbb{M}'_1$ for some G'_h and $\mathscr{P}'_h \subseteq \mathscr{P}_h$ and for all $h \in H$. We conclude $G \xrightarrow{p\lambda q} r \to s : \{\lambda'_h.G'_h\}_{h \in H}$ using Rule [ICOMM] and $r \to s : \{\lambda'_h.G'_h\}_{h \in H} \vdash_{\mathscr{P}'} \mathbb{M}'$ with $\mathscr{P}' = \bigcup_{h \in H} \mathscr{P}'_h$ using Rule [COMM].

Rule [WEAK]. In this case $\mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $\mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2$ and $G \vdash_{\mathscr{P}_1} \mathbb{M}_1$. From $\{p,q\} \subseteq prt(G)$, by Lemma 3.3 we get $\{p,q\} \subseteq prt(\mathbb{M}_1)$ which implies $\mathbb{M}_1 \xrightarrow{p\lambda q} \mathbb{M}'_1$. By coinduction we get $G \xrightarrow{p\lambda q} G'$ and $G' \vdash_{\mathscr{P}'_1} \mathbb{M}'_1$ with $\mathscr{P}'_1 \subseteq \mathscr{P}_1$. We conclude using Rule [WEAK], since by construction $\mathbb{M}' \equiv \mathbb{M}'_1 \parallel \mathbb{M}_2$. $\square$

We note that Subject Reduction, as formulated in previous theorem, fails if we allow unbounded global types. Let $\mathsf{G} = \mathsf{p} \to \mathsf{q}{:}\{\lambda_1.\mathsf{r} \to \mathsf{s}{:}\lambda, \lambda_2.\mathsf{G}\}$ and $\mathbb{M} \equiv \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{r}[\mathsf{s}!\lambda] \parallel \mathsf{s}[\mathsf{r}?\lambda]$ and $P = \mathsf{q}!\{\lambda_1, \lambda_2.P\}$ and $Q = \mathsf{p}?\{\lambda_1, \lambda_2.Q\}$. Then we have $\mathsf{G} \vdash_\emptyset \mathbb{M}$ and $\mathbb{M} \xrightarrow{\mathsf{r}\lambda\mathsf{s}} \mathsf{p}[P] \parallel \mathsf{q}[Q]$, but there is no transition labelled $\mathsf{r}\lambda\mathsf{s}$ starting from G. Note that the session $\mathbb{M}$ can be typed, still with the $\emptyset$ subscript, by the bounded global type $\mathsf{G}' = \mathsf{r} \to \mathsf{s}{:}\lambda.\mathsf{p} \to \mathsf{q}{:}\{\lambda_1, \lambda_2.\mathsf{G}'\}$.

Session Fidelity ensures that the communications in a session typed by a global type proceed at least as prescribed by the global type.

Theorem 3.7 (Session Fidelity) *Let* $\mathsf{G} \vdash_\mathscr{P} \mathbb{M}$ *and* $\mathsf{G} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathsf{G}'$. *Then* $\mathbb{M} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathbb{M}'$ *and* $\mathsf{G}' \vdash_{\mathscr{P}'} \mathbb{M}'$ *with* $\mathscr{P}' \subseteq \mathscr{P}$.

Proof. The proof is by coinduction on the derivation of $\mathsf{G} \vdash_\mathscr{P} \mathbb{M}$ and by cases on the last applied rule.

Rule [COMM]. The proof is by induction on the number t of transition rules used to derive $\mathsf{G} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathsf{G}'$. *Case* $t = 1$. Then $\mathsf{G} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathsf{G}'$ is the Axiom [ECOMM] and $\mathsf{G} = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\mathsf{G}_i\}_{i \in I}$, where $\lambda = \lambda_l$ and $\mathsf{G}' = \mathsf{G}_l$ with $l \in I$. We get $\mathbb{M} \equiv \mathsf{p}[\mathsf{q}!\{\lambda_i.P_i\}_{i \in I}] \parallel \mathsf{q}[\mathsf{p}?\{\lambda_j.Q_j\}_{j \in J}] \parallel \mathbb{M}_0$ with $I \subseteq J$ and $\mathsf{G}_i \vdash_{\mathscr{P}_i} \mathsf{p}[P_i] \parallel \mathsf{q}[Q_i] \parallel \mathbb{M}_0$ for all $i \in I$ with $\mathscr{P} = \bigcup_{i \in I} \mathscr{P}_i$. Then we conclude $\mathbb{M} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathsf{p}[P_l] \parallel \mathsf{q}[Q_l] \parallel \mathbb{M}_0$ by Rule [COMM-T] and $\mathsf{G}_l \vdash_{\mathscr{P}_l} \mathsf{p}[P_l] \parallel \mathsf{q}[Q_l] \parallel \mathbb{M}_0$.
Case $t > 1$. Then $\mathsf{G} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathsf{G}'$ is the conclusion of Rule [ICOMM]. Moreover $\mathsf{G} = \mathsf{r} \to \mathsf{s} : \{\lambda_h'.\mathsf{G}_h\}_{h \in H}$ and $\mathsf{G}' = \mathsf{r} \to \mathsf{s} : \{\lambda_h'.\mathsf{G}_h'\}_{h \in H}$ and $\mathsf{G}_h \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathsf{G}_h'$ for all $h \in H$ and $\{\mathsf{p},\mathsf{q}\} \cap \{\mathsf{r},\mathsf{s}\} = \emptyset$. We get

$$\mathbb{M} \equiv \mathsf{r}[\mathsf{s}!\{\lambda_h'.R_h\}_{h \in H}] \parallel \mathsf{s}[\mathsf{r}?\{\lambda_k'.S_k\}_{k \in K}] \parallel \mathbb{M}_1$$

with $H \subseteq K$ and $\mathsf{G}_h \vdash_{\mathscr{P}_h} \mathsf{r}[R_h] \parallel \mathsf{s}[S_h] \parallel \mathbb{M}_1$ for all $h \in H$ with $\mathscr{P} = \bigcup_{h \in H} \mathscr{P}_h$. By induction $\mathsf{r}[R_h] \parallel \mathsf{s}[S_h] \parallel \mathbb{M}_1 \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathbb{M}_h'$ and $\mathsf{G}_h' \vdash_{\mathscr{P}_h'} \mathbb{M}_h'$ with $\mathscr{P}_h' \subseteq \mathscr{P}_h$ for all $h \in H$. The condition $\{\mathsf{p},\mathsf{q}\} \cap \{\mathsf{r},\mathsf{s}\} = \emptyset$ ensures that the reduction $\mathsf{r}[R_h] \parallel \mathsf{s}[S_h] \parallel \mathbb{M}_1 \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathbb{M}_h'$ does not modify the processes of participants r and s. Moreover the processes of participants p and q are the same in $\mathbb{M}_h'$ for all $h \in H$. This implies $\mathbb{M}_h' \equiv \mathsf{r}[R_h] \parallel \mathsf{s}[S_h] \parallel \mathbb{M}''$ for all $h \in H$ and some $\mathbb{M}''$. We conclude $\mathbb{M} \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathbb{M}'$ where $\mathbb{M}' = \mathsf{r}[\mathsf{s}!\{\lambda_h'.R_h\}_{h \in H}] \parallel \mathsf{s}[\mathsf{r}?\{\lambda_k'.S_k\}_{k \in K}] \parallel \mathbb{M}''$ using Rule [COMM-T] and $\mathsf{G}' \vdash_{\mathscr{P}'} \mathbb{M}'$ with $\mathscr{P}' = \bigcup_{h \in H} \mathscr{P}_h'$ using Rule [COMM].

Rule [WEAK]. In this case $\mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $\mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2$ and $\mathsf{G} \vdash_{\mathscr{P}_1} \mathbb{M}_1$. By coinduction $\mathbb{M}_1 \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathbb{M}_1'$ and $\mathsf{G}' \vdash_{\mathscr{P}_1'} \mathbb{M}_1'$ with $\mathscr{P}_1' \subseteq \mathscr{P}_1$. Then $\mathbb{M}_1 \parallel \mathbb{M}_2 \xrightarrow{\mathsf{p}\lambda\mathsf{q}} \mathbb{M}_1' \parallel \mathbb{M}_2$ and $\mathsf{G}' \vdash_{\mathscr{P}_1' \cup \mathscr{P}_2} \mathbb{M}_1' \parallel \mathbb{M}_2$ using Rule [WEAK]. $\square$

We can show that typability ensures $\mathscr{P}$-excluded Lock-freedom. This follows from Subject Reduction and Session Fidelity thanks to the boundedness condition.

Theorem 3.8 ($\mathscr{P}$-excluded Lock-freedom) *If* $\mathsf{G} \vdash_\mathscr{P} \mathbb{M}$, *then* $\mathbb{M}$ *is* $\mathscr{P}$-*excluded lock-free.*

Proof. By Subject Reduction it is enough to prove that in well-typed sessions no active participant not belonging to $\mathscr{P}$ is prevented to progress. So, let $\mathsf{p} \in \mathsf{prt}(\mathbb{M})$ such that $\mathsf{p} \notin \mathscr{P}$. By Lemma 3.3 we have that $\mathsf{p} \in \mathsf{prt}(\mathbb{M})$ and $\mathsf{p} \notin \mathscr{P}$ imply $\mathsf{p} \in \mathsf{prt}(\mathsf{G})$. We proceed now by induction on $d = depth(\mathsf{G}, \mathsf{p})$.
If $d = 1$ then either $\mathsf{G} = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\mathsf{G}_i\}_{i \in I}$ or $\mathsf{G} = \mathsf{q} \to \mathsf{p} : \{\lambda_i.\mathsf{G}_i\}_{i \in I}$ and $\mathsf{G} \xrightarrow{\Lambda} \mathsf{G}'$ with $\mathsf{p} \in \mathsf{prt}(\Lambda)$ by Axiom [ECOMM]. Then $\mathbb{M} \xrightarrow{\Lambda} \mathbb{M}'$ by Theorem 3.7.
If $d > 1$ then $\mathsf{G} = \mathsf{q} \to \mathsf{r} : \{\lambda_i.\mathsf{G}_i\}_{i \in I}$ with $\mathsf{p} \notin \{\mathsf{q},\mathsf{r}\}$ and $\mathsf{G} \xrightarrow{\mathsf{q}\lambda_i\mathsf{r}} \mathsf{G}_i$ for all $i \in I$ by Axiom [ECOMM]. Induction applies since $depth(\mathsf{G}, \mathsf{p}) > depth(\mathsf{G}_i, \mathsf{p})$ for all $i \in I$. Then, for all $i \in I$, we get $\mathsf{G}_i \xrightarrow{\sigma_i \cdot \Lambda_i} \mathsf{G}_i'$ for

some σ_i, Λ_i with $\mathsf{p} \in \mathsf{prt}(\Lambda_i)$. This implies $\mathsf{G} \xrightarrow{\sigma_i'} \mathsf{G}_i'$ where $\sigma_i' = \mathsf{q}\lambda_i\mathsf{r} \cdot \sigma_i \cdot \Lambda_i$ for all $i \in I$. We conclude $\mathbb{M} \xrightarrow{\sigma_i'} \mathbb{M}_i'$ for all $i \in I$ by Theorem 3.7. $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\square$

The following example shows that partial typing allows to type sessions which require unbounded global types in standard type systems [17].

Example 3.9 (Buyer-Seller-Carrier) Let us consider the following session (from [11, Sect.1])

$$\mathbb{M} = \mathsf{b}[B] \parallel \mathsf{s}[S] \parallel \mathsf{c}[\mathsf{s?}\textsc{ship}]$$

where $B = \mathsf{s!}\{\textsc{add}.B, \textsc{pay}\}$ and $S = \mathsf{b?}\{\textsc{add}.S, \textsc{pay}.\mathsf{c!}\textsc{ship}\}$

Such a session implements a system where a buyer can keep on adding goods - sold by a seller - in his shopping cart an unbounded number of times, until he decides to buy the shopping cart's content. In the latter case, the seller informs the carrier for the shipment. Session $\mathbb{M}$ is obviously non lock-free, since participant c would not be able to progress in case s be a seriously disturbed shopaholic who never stop adding goods in his cart. In fact $\mathbb{M}$ cannot be typed in $\vdash_\emptyset$ (which corresponds to the type system of [3]).

In this scenario participant b is a client, whereas s and c are part of the service used by b. It is hence natural to look at $\mathbb{M}$ with a bias towards the client, the one whose good property have to be ensured. As a matter of fact it is possible to ensure the lock freedom of b, namely the $\{\mathsf{s},\mathsf{c}\}$-excluded lock freedom of $\mathbb{M}$, by deriving

$$\mathsf{b} \rightarrow \mathsf{s}{:}\{\textsc{add}. \mathsf{G}, \textsc{buy}\} \vdash_{\{\mathsf{s},\mathsf{c}\}} \mathbb{M}$$

as follows

$$\mathscr{D} \quad = \quad \cfrac{\mathscr{D} \qquad\qquad \cfrac{\cfrac{\overline{\overline{\mathsf{End} \vdash_\emptyset \mathsf{b}[\mathbf{0}]}} \text{ [End]}}{\mathsf{End} \vdash_{\{\mathsf{s},\mathsf{c}\}} \mathsf{b}[\mathbf{0}] \parallel \mathsf{s}[\mathsf{c!}\textsc{ship}] \parallel \mathsf{c}[\mathsf{s?}\textsc{ship}]} \text{ [S-\textsc{weak}]}}{\mathsf{b} \rightarrow \mathsf{s}{:}\{\textsc{add}. \mathsf{G}, \textsc{buy}\} \vdash_{\{\mathsf{s},\mathsf{c}\}} \mathsf{b}[\mathsf{s!}\{\textsc{add}.B, \textsc{pay}\}] \parallel \mathsf{s}[\mathsf{b?}\{\textsc{add}.S, \textsc{pay}.\mathsf{c!}\textsc{ship}\}] \parallel \mathsf{c}[\mathsf{s?}\textsc{ship}]}$$

where $\mathsf{G} = \mathsf{b} \rightarrow \mathsf{s}{:}\{\textsc{add}. \mathsf{G}, \textsc{buy}\}$. Note that $\mathsf{b} \rightarrow \mathsf{s}{:}\{\textsc{add}. \mathsf{G}, \textsc{buy}\}$ is a bounded global type. $\qquad\diamond$

4 Type Inference

In our type system each session can be trivially typed by the End type just applying Axiom [End] and Rule [Weak]:

$$\mathsf{End} \vdash_{\mathsf{prt}(\mathbb{M})} \mathbb{M}$$

Clearly, this typing does not provide any information on $\mathbb{M}$. We are interested in more informative typings, if any. In this section, we will describe an algorithm to infer global types and sets of participants from sessions, proving also its soundness and completeness with respect to our type system. In particular, the algorithm applied to a session $\mathbb{M}$ returns all and only those global types which can be assigned to $\mathbb{M}$ with derivations indexed by suitable sets of participants. Note that, since derivations indexed by the same or different sets of participants can assign different global types to a session, the algorithm needs to be non-deterministic in order to be complete.

The first step towards defining such an algorithm is the introduction of a finite representation for global types. Since global types are regular terms, they can be represented, by results in [1, 12], as finite systems of regular syntactic equations formally defined below.

We begin by defining a *global type pattern* as a finite term generated by the following grammar:

$$\mathbb{G} ::= \mathtt{End} \mid \mathtt{p} \to \mathtt{q} : \{\lambda_i.\mathbb{G}_i\}_{i\in I} \mid X$$

where X is a type variable taken from a countably infinite set. We denote by $\mathsf{vars}(\mathbb{G})$ the set of type variables occurring in $\mathbb{G}$. We also need to compute sets of participants, so we define p-set patterns by:

$$\mathbb{P} ::= \mathscr{P} \mid x \mid \mathbb{P} \cup \mathbb{P}$$

where x is a p-set variable taken from a countably infinite set and $\mathscr{P}$ can be any finite set of participants. We denote by $\mathsf{vars}(\mathbb{P})$ the set of p-set variables occurring in $\mathbb{P}$.

We use χ to range over type and p-set variables.

A *substitution* θ is a finite partial map from type variables to global types and from p-set variables to sets of participants. We denote by $\theta + \sigma$ the union of two substitutions such that $\theta(\chi) = \sigma(\chi)$, for all $\chi \in \mathsf{dom}(\theta) \cap \mathsf{dom}(\sigma)$, and by $\mathbb{G}\theta$ (resp. $\mathbb{P}\theta$) the application of θ to $\mathbb{G}$ (resp. $\mathbb{P}$). We define $\theta \preceq \sigma$ if $\mathsf{dom}(\theta) \subseteq \mathsf{dom}(\sigma)$ and $\theta(\chi) = \sigma(\chi)$, for all $\chi \in \mathsf{dom}(\theta)$. Note that, if $\mathsf{vars}(\mathbb{G}) \subseteq \mathsf{dom}(\theta)$, then $\mathbb{G}\theta$ is a global type and if $\mathsf{vars}(\mathbb{P}) \subseteq \mathsf{dom}(\theta)$, then $\mathbb{P}\theta$ is a set of participants.

A *type equation* has shape $X = \mathbb{G}$ and a *(regular) system of type equations* $\mathscr{E}$ is a finite set of equations such that $X = \mathbb{G}_1$ and $X = \mathbb{G}_2 \in \mathscr{E}$ imply $\mathbb{G}_1 = \mathbb{G}_2$. We denote by $\mathsf{dom}(\mathscr{E})$ the set $\{X \mid X = \mathbb{G} \in \mathscr{E}\}$ and by $\mathsf{vars}(\mathscr{E})$ the set $\bigcup\{\mathsf{vars}(\mathbb{G}) \cup \{X\} \mid X = \mathbb{G} \in \mathscr{E}\}$. A *solution* of a system $\mathscr{E}$ is a substitution θ such that $\mathsf{vars}(\mathscr{E}) \subseteq \mathsf{dom}(\theta)$ and, for all $X = \mathbb{G} \in \mathscr{E}$, $\theta(X) = \mathbb{G}\theta$ holds and $\theta(X)$ is bounded. We denote by $\mathsf{sol}(\mathscr{E})$ the set of all solutions of $\mathscr{E}$.

A *p-set equation* has shape $x = \mathbb{P}$. We use E to range over regular systems of p-set equations, which are defined similarly to regular systems of type equations. Also $\mathsf{dom}(E)$, $\mathsf{vars}(E)$ and $\mathsf{sol}(E)$ have the same meanings as for systems of type equations.

A *p-condition* has shape $(\mathsf{prt}(X) \cup x) \setminus \{\mathtt{p},\mathtt{q}\} \doteq \mathscr{P}$ and we let $\mathscr{C}$ range over sets of p-conditions with pairwise distinct type and p-set variables. A substitution θ *agrees* with
- $(\mathsf{prt}(X) \cup x) \setminus \{\mathtt{p},\mathtt{q}\} \doteq \mathscr{P}$ if $(\mathsf{prt}(\theta(X)) \cup \theta(x)) \setminus \{\mathtt{p},\mathtt{q}\} = \mathscr{P}$;
- $\mathscr{C}$, notation $\theta \propto \mathscr{C}$, if θ agrees with all p-conditions in $\mathscr{C}$.

We define $\mathsf{sol}(\mathscr{E},E,\mathscr{C})$ as the set of solutions of $\mathscr{E}$ and E which agree with $\mathscr{C}$, i.e. $\mathsf{sol}(\mathscr{E},E,\mathscr{C}) = \{\theta \in \mathsf{sol}(\mathscr{E}) \cap \mathsf{sol}(E) \mid \theta \propto \mathscr{C}\}$ and note that $\mathscr{E}_1 \subseteq \mathscr{E}_2$, $E_1 \subseteq E_2$, $\mathscr{C}_1 \subseteq \mathscr{C}_2$ imply $\mathsf{sol}(\mathscr{E}_2,E_2,\mathscr{C}_2) \subseteq \mathsf{sol}(\mathscr{E}_1,E_1,\mathscr{C}_1)$.

The algorithm follows essentially the structure of coSLD resolution of coinductive logic programming [29, 30, 31, 2], namely the extension of standard SLD resolution capable to deal with regular terms and coinductive predicates. A *goal* is a triple $(X,\mathbb{M},x)$ of a type variable X, a session $\mathbb{M}$ and a p-set variable x. The algorithm takes a goal $(X,\mathbb{M},x)$ as input, and returns a system of type equations $\mathscr{E}$ and a system of p-set equations E and a set of p-condition $\mathscr{C}$. A solution for the variable X in $\mathscr{E}$ is a global type for the session $\mathbb{M}$ in a derivation indexed by a solution for the variable x in E which satisfies the p-conditions in $\mathscr{C}$. The key idea, borrowed from coinductive logic programming, is to keep track of already encountered goals in order to detect cycles and so avoiding non-termination.

The inference judgements have the following shape: $\mathscr{S} \vdash (X,\mathbb{M},x) \Rightarrow (\mathscr{E},E,\mathscr{C})$, where $\mathscr{S}$ is a set of goals, all with different variables which are all different from X and x. Rules defining the inference algorithm are reported in Figure 2.

For a terminated session the algorithm returns just the two equations $X = \mathtt{End}$ and $x = \emptyset$ and the empty set of conditions (Axiom [A-END]).

In Rule [A-COMM] the algorithm nondeterministically selects one of the matching pairs of processes: $P = \mathtt{q}!\{\lambda_i.P_i\}_{i\in I}$ and $Q = \mathtt{p}?\{\lambda_j.Q_j\}_{j\in J}$, with $I \subseteq J$, i.e. two participants willing to communicate such that the output process can freely choose the message. The algorithm is then recursively applied, for each $i \in I$, to the session where the processes of $\mathtt{p}$ and $\mathtt{q}$ are, respectively, P_i and Q_i. In each call the goal $(X,\mathtt{p}[P] \parallel \mathtt{q}[Q] \parallel \mathbb{M},x)$ is added to the set of goals. At the end of the recursive calls the algorithm

$$[\text{A-End}] \; \frac{}{\mathscr{S} \vdash (X, \mathsf{p}[\mathbf{0}], x) \Rightarrow (\{X \doteq \mathtt{End}\}, \{x \doteq \emptyset\}, \emptyset)}$$

$$[\text{A-Cycle}] \; \frac{}{\mathscr{S}, (Y, \mathbb{M}, y) \vdash (X, \mathbb{M}, x) \Rightarrow (\{X \doteq Y\}, \{x \doteq y\}, \emptyset)}$$

$$[\text{A-Comm}] \; \frac{\mathscr{S}' \vdash (Y_i, \mathsf{p}[P_i] \parallel \mathsf{q}[Q_i] \parallel \mathbb{M}, y_i) \Rightarrow (\mathscr{E}_i, E_i, \mathscr{C}_i) \;\; \forall i \in I}{\mathscr{S} \vdash (X, \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})} \quad \begin{array}{l} \mathscr{S}' = \mathscr{S}, (X, \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathbb{M}, x) \\ P = \mathsf{q}!\{\lambda_i.P_i\}_{i \in I} \\ Q = \mathsf{p}?\{\lambda_j.Q_j\}_{j \in J} \quad\quad I \subseteq J \\ Y_i, y_i \text{ fresh } \forall i \in I \\ \mathscr{E} = \{X \doteq \mathsf{p} \to \mathsf{q} : \{\lambda_i.Y_i\}_{i \in I}\} \cup \bigcup_{i \in I} \mathscr{E}_i \\ E = \{x \doteq \bigcup_{i \in I} y_i\} \cup \bigcup_{i \in I} E_i \\ \mathscr{C} = \{(\mathsf{prt}(Y_i) \cup y_i) \setminus \{\mathsf{p}, \mathsf{q}\} \doteq \mathsf{prt}(\mathbb{M}) \mid \forall i \in I\} \\ \quad\quad \cup \bigcup_{i \in I} \mathscr{C}_i \end{array}$$

$$[\text{A-Weak}] \; \frac{\mathscr{S}, (X, \mathbb{M}_1 \parallel \mathbb{M}_2, x) \vdash (Y, \mathbb{M}_1, y) \Rightarrow (\mathscr{E}_1, E_1, \mathscr{C})}{\mathscr{S} \vdash (X, \mathbb{M}_1 \parallel \mathbb{M}_2, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})} \quad \begin{array}{l} Y, y \text{ fresh} \\ \mathscr{P} = \mathsf{prt}(\mathbb{M}_2) \neq \emptyset \\ \mathscr{E} = \{X \doteq Y\} \cup \mathscr{E}_1 \\ E = \{x \doteq y \cup \mathscr{P}\} \cup E_1 \end{array}$$

Figure 2: Rules of the inference algorithm.

collects all the resulting equations plus another two for the current variables. Note that variables for the goals in the premises are fresh. This is important to ensure that the sets of equations $\mathscr{E}$ and E in the conclusion are indeed regular systems of equations (there is at most one equation for each variable). The new p-condition ensures that the resulting global type associated to X and the resulting set of participants associated to x satisfy the conditions on participants required by Rule [Comm] in Definition 3.1.

In Rule [A-Weak] the algorithm nondeterministically partitions the input session into two subsessions $\mathbb{M}_1$ and $\mathbb{M}_2$ and then it is recursively called on the former. After the recursive call it simply adds the same participants to the session (together with their processes) and to the set of ignored participants by means of the equation $x \doteq y \cup \mathscr{P}$.

Finally, Axiom [A-Cycle] detects cycles: if the session in the current goal appears also in the set $\mathscr{S}$, the algorithm can stop, returning just two equations unifying the type and p-set variables associated with the session together with the empty set of conditions.

Example 4.1 (Inference for the social media) Figure 3 gives a type inference where:
- the processes P, Q, U, P_1, Q_1, U_1, P_2, Q_2 are defined as in Example 3.2 and $U_2 = \mathsf{q}!\text{DND}.U$, $U_3 = \mathsf{q}!\text{GRTD}.U$;
- the goals are

$$\mathscr{S}_1 = \{(X, \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{u}[U], x)\} \quad\quad \mathscr{S}_2 = \mathscr{S}_1 \cup \{(Y_1, \mathsf{p}[P_1] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U], y_1)\}$$

$$\mathscr{S}_3 = \mathscr{S}_4 = \mathscr{S}_2 \cup \{(Y_2, \mathsf{p}[\mathsf{u}?\{\text{DND}.P, \text{GRTD}.P_2\}] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_1], y_2)\}$$

$$\mathscr{S}_5 = \mathscr{S}_3 \cup \{(Y_3, \mathsf{p}[P] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_2], y_3)\} \quad\quad \mathscr{S}_6 = \mathscr{S}_4 \cup \{(Y_4, \mathsf{p}[P_2] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_3], y_4)\}$$

$$\mathscr{S}_7 = \mathscr{S}_6 \cup \{(Y_6, \mathsf{p}[P_2] \parallel \mathsf{q}[Q_2] \parallel \mathsf{u}[U], y_6)\} \quad\quad \mathscr{S}_8 = \mathscr{S}_7 \cup \{(Y_7, \mathsf{p}[P_2] \parallel \mathsf{q}[Q_2], y_7)\}$$

- the systems of type equations are

$$\mathscr{E} = \{X \doteq \mathsf{q} \to \mathsf{p}:\text{HELLO}.Y_1\} \cup \mathscr{E}_1 \quad\quad \mathscr{E}_1 = \{Y_1 \doteq \mathsf{p} \to \mathsf{u}:\text{REQ}.Y_2\} \cup \mathscr{E}_2$$

$$\frac{\mathscr{S}_8 \vdash (Y_8, \mathsf{p}[\mathbf{0}] \parallel \mathsf{q}[\mathbf{0}], y_8) \Rightarrow (\mathscr{E}_8, E_8, \mathscr{C}_7)}{\mathscr{S}_7 \vdash (Y_7, \mathsf{p}[P_2] \parallel \mathsf{q}[Q_2], y_7) \Rightarrow (\mathscr{E}_7, E_7, \mathscr{C}_6)}$$

$$\frac{\mathscr{S}_5 \vdash (Y_5, \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{u}[U], y_5) \Rightarrow (\mathscr{E}_5, E_5, \mathscr{C}_5)}{\mathscr{S}_3 \vdash (Y_3, \mathsf{p}[P] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_2], y_3) \Rightarrow (\mathscr{E}_3, E_3, \mathscr{C}_3)} \qquad \frac{\mathscr{S}_6 \vdash (Y_6, \mathsf{p}[P_2] \parallel \mathsf{q}[Q_2] \parallel \mathsf{u}[U], y_6) \Rightarrow (\mathscr{E}_6, E_6, \mathscr{C}_6)}{\mathscr{S}_4 \vdash (Y_4, \mathsf{p}[P_2] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_3], y_4) \Rightarrow (\mathscr{E}_4, E_4, \mathscr{C}_4)}$$

$$\frac{\mathscr{S}_2 \vdash (Y_2, \mathsf{p}[\mathsf{u}?\{\text{DND}.P, \text{GRTD}.P_2\}] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U_1], y_2) \Rightarrow (\mathscr{E}_2, E_2, \mathscr{C}_2)}{\mathscr{S}_1 \vdash (Y_1, \mathsf{p}[P_1] \parallel \mathsf{q}[Q_1] \parallel \mathsf{u}[U], y_1) \Rightarrow (\mathscr{E}_1, E_1, \mathscr{C}_1)}$$

$$\frac{}{\vdash (X, \mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathsf{u}[U], x) \Rightarrow (\mathscr{E}, E, \mathscr{C})}$$

Figure 3: A type inference for the social media.

$$\mathscr{E}_2 = \{Y_2 = \mathsf{u} \to \mathsf{p}:\{\text{DND}.Y_3, \text{GRTD}.Y_4\}\} \cup \mathscr{E}_3 \cup \mathscr{E}_4 \qquad \mathscr{E}_3 = \{Y_3 = \mathsf{u} \to \mathsf{q}:\text{DND}.Y_5\} \cup \mathscr{E}_5$$

$$\mathscr{E}_4 = \{Y_4 = \mathsf{u} \to \mathsf{q}:\text{GRTD}.Y_6\} \cup \mathscr{E}_6 \qquad \mathscr{E}_5 = \{Y_5 = X\}$$

$$\mathscr{E}_6 = \{Y_6 = Y_7\} \cup \mathscr{E}_7 \qquad \mathscr{E}_7 = \{Y_7 = \mathsf{p} \to \mathsf{q}:\text{HELLO}.Y_8\} \cup \mathscr{E}_8 \qquad \mathscr{E}_8 = \{Y_8 = \texttt{End}\}$$

- the systems of p-set equations are

$$E = \{x = y_1\} \cup E_1 \qquad E_1 = \{y_1 = y_2\} \cup E_2 \qquad E_2 = \{y_2 = y_3 \cup y_4\} \cup E_3 \cup E_4$$

$$E_3 = \{y_3 = y_5\} \cup E_5 \qquad E_4 = \{y_4 = y_6\} \cup E_6 \qquad E_5 = \{y_5 = x\}$$

$$E_6 = \{y_6 = y_7 \cup \{\mathsf{u}\}\} \cup E_7 \qquad E_7 = \{y_7 = y_8\} \cup E_8 \qquad E_8 = \{y_8 = \emptyset\}$$

- the sets of p-conditions are

$$\mathscr{C} = \{(\mathsf{prt}(Y_1) \cup y_1) \setminus \{\mathsf{p}, \mathsf{q}\} = \{\mathsf{u}\}\} \cup \mathscr{C}_1 \qquad \mathscr{C}_1 = \{(\mathsf{prt}(Y_2) \cup y_2) \setminus \{\mathsf{p}, \mathsf{u}\} = \{\mathsf{q}\}\} \cup \mathscr{C}_2$$

$$\mathscr{C}_2 = \{(\mathsf{prt}(Y_3) \cup y_3) \setminus \{\mathsf{u}, \mathsf{p}\} = \{\mathsf{q}\}, (\mathsf{prt}(Y_4) \cup y_4) \setminus \{\mathsf{u}, \mathsf{p}\} = \{\mathsf{q}\}\} \cup \mathscr{C}_3 \cup \mathscr{C}_4$$

$$\mathscr{C}_3 = \{(\mathsf{prt}(Y_5) \cup y_5) \setminus \{\mathsf{u}, \mathsf{q}\} = \{\mathsf{p}\}\} \cup \mathscr{C}_5 \qquad \mathscr{C}_4 = \{(\mathsf{prt}(Y_6) \cup y_6) \setminus \{\mathsf{u}, \mathsf{q}\} = \{\mathsf{p}\}\} \cup \mathscr{C}_6 \qquad \mathscr{C}_5 = \emptyset$$

$$\mathscr{C}_6 = \{y_6 = y_7 \cup \{\mathsf{u}\}\} \cup \mathscr{C}_7 \qquad \mathscr{C}_7 = \{(\mathsf{prt}(Y_8) \cup y_8) \setminus \{\mathsf{p}, \mathsf{q}\} = \emptyset\} \cup \mathscr{C}_8 \qquad \mathscr{C}_8 = \emptyset$$

One can easily verify that a solution of both systems of equations $\mathscr{E}$ and E satisfying the p-conditions (i.e. which agrees with $\mathscr{C}$) is $X = \mathsf{G}$ and $x = \{\mathsf{u}\}$, where G is the global type defined in Example 2.12 and derived for this session in Figure 1. $\Diamond$

Let $\mathscr{E}$, E be two systems of type and p-set equations, $\mathscr{C}$ a set of p-conditions and $\mathscr{S}$ a set of goals. A solution $\theta \in \mathsf{sol}(\mathscr{E}, E, \mathscr{C})$ *agrees* with $\mathscr{S}$ if $(X, \mathbb{M}, x) \in \mathscr{S}$ implies $\mathsf{prt}(\theta(X)) \cup \theta(x) = \mathsf{prt}(\mathbb{M})$ for all $X \in \mathsf{vars}(\mathscr{E})$ and all $x \in \mathsf{vars}(E)$. We denote by $\mathsf{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$ the set of all solutions in $\mathsf{sol}(\mathscr{E}, E, \mathscr{C})$ agreeing with $\mathscr{S}$. We say that a system of type equations $\mathscr{E}$ is *guarded* if $X = Y$ and $Y = \mathbb{G}$ in $\mathscr{E}$ imply that $\mathbb{G}$ is not a variable. Moreover, $\mathscr{E}$ is $\mathscr{S}$-*closed* if it is guarded and $\mathsf{dom}(\mathscr{E}) \cap \mathsf{vars}(\mathscr{S}) = \emptyset$ and $\mathsf{vars}(\mathscr{E}) \setminus \mathsf{dom}(\mathscr{E}) \subseteq \mathsf{vars}(\mathscr{S})$. We define similarly when a set of p-set equations E is guarded and $\mathscr{S}$-closed.

Toward proving properties of the inference algorithm, we check a couple of auxiliary lemmas. As usual $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ means that this judgment belongs to a derivation in the system of Figure 2 having a judgment with an empty sets of goals as conclusion (namely it represents the result of a recursive call during the execution of our algorithm).

$$[\iota\text{-}\textsc{End}] \ \frac{}{\mathcal{N} \vdash^\iota_\emptyset \mathsf{p[0]} : \mathtt{End}} \qquad\qquad [\iota\text{-}\textsc{Cycle}] \ \frac{}{\mathcal{N}, (\mathsf{G}, \mathbb{M}, \mathscr{P}) \vdash^\iota_{\mathscr{P}} \mathbb{M} : \mathsf{G}}$$

$$[\iota\text{-}\textsc{Comm}] \ \frac{\begin{array}{c} \mathcal{N}, (\mathsf{G}, \mathbb{M}, \mathscr{P}) \vdash^\iota_{\mathscr{P}_i} \mathsf{p}[P_i] \parallel \mathsf{q}[Q_i] \parallel \mathbb{M}' : \mathsf{G}_i \\ (\mathsf{prt}(\mathsf{G}_i) \cup \mathscr{P}_i) \setminus \{\mathsf{p}, \mathsf{q}\} = \mathsf{prt}(\mathbb{M}') \ \ \forall i \in I \end{array}}{\mathcal{N} \vdash^\iota_{\mathscr{P}} \mathbb{M} : \mathsf{G}} \qquad \begin{array}{c} \mathbb{M} \equiv \mathsf{p}[\mathsf{q}!\{\lambda_i.P_i\}_{i \in I}] \parallel \mathsf{q}[\mathsf{p}?\{\lambda_j.Q_j\}_{j \in J}] \parallel \mathbb{M}' \\ \mathsf{G} = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\mathsf{G}_i\}_{i \in I} \quad \mathsf{G} \text{ is bounded} \\ \mathscr{P} = \cup_{i \in I} \mathscr{P}_i \quad I \subseteq J \end{array}$$

$$[\iota\text{-}\textsc{Weak}] \ \frac{\mathcal{N}, (\mathsf{G}, \mathbb{M}_1 \parallel \mathbb{M}_2, \mathscr{P}_1 \cup \mathscr{P}_2) \vdash^\iota_{\mathscr{P}_1} \mathbb{M}_1 : \mathsf{G}}{\mathcal{N} \vdash^\iota_{\mathscr{P}_1 \cup \mathscr{P}_2} \mathbb{M}_1 \parallel \mathbb{M}_2 : \mathsf{G}} \quad \mathscr{P}_2 = \mathsf{prt}(\mathbb{M}_2) \neq \emptyset$$

Figure 4: Inductive typing rules for sessions.

Lemma 4.2 *If $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$, then $\mathscr{E}$ and E are $\mathscr{S}$-closed.*

Proof. By a straightforward induction on the derivation of $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$. $\qquad\square$

Lemma 4.3 *If $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ and $\theta \in \mathsf{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$, then $\mathsf{prt}(\theta(X)) \cup \theta(x) = \mathsf{prt}(\mathbb{M})$.*

Proof. By induction on the derivation of $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$. The only interesting case is when Rule [A-Comm] is applied. From $\theta \propto \mathscr{C}$ we get $(\mathsf{prt}(\theta(Y_i)) \cup \theta(y_i)) \setminus \{\mathsf{p}, \mathsf{q}\} = \mathsf{prt}(\mathbb{M})$ for all $i \in I$, which imply $\mathsf{prt}(\theta(X)) \cup \theta(x) = \mathsf{prt}(\mathsf{p}[P] \parallel \mathsf{q}[Q] \parallel \mathbb{M})$ since $X \Equiv \mathsf{p} \to \mathsf{q} : \{\lambda_i.Y_i\}_{i \in I} \in \mathscr{E}$ and $x \Equiv \bigcup_{i \in I} y_i \in E$. $\square$

To show soundness and completeness of our inference algorithm, it is handy to formulate an inductive version of our typing rules, see Figure 4, where $\mathcal{N}$ ranges over sets of triples $(\mathsf{G}, \mathbb{M}, \mathscr{P})$. We can give an inductive formulation since all infinite derivations using the typing rules of Definition 3.1 are regular, i.e. the number of different subtrees of a derivation for a judgement $\mathsf{G} \vdash_{\mathscr{P}} \mathbb{M}$ is finite. In fact, it is bounded by the product of the number of different subterms of G and the number of different subsessions of $\mathbb{M}$, which are both finite as G and (processes in) $\mathbb{M}$ are regular. Applying the standard transformation according to [27, Section 21.9] from a coinductive to an inductive formulation we get the typing rules shown in Figure 4.

In the following two lemmas we relate inference and inductive derivability.

Lemma 4.4 *If $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$, then $\mathscr{S}\theta \vdash^\iota_{\theta(x)} \mathbb{M} : \theta(X)$ for all $\theta \in \mathsf{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$ such that* $\mathsf{vars}(\mathscr{S}) \subseteq \mathsf{dom}(\theta)$.

Proof. By induction on the derivation of $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$.

Axiom [A-End] We have $\mathscr{E} = \{X \Equiv \mathtt{End}\}$, $E = \{x \Equiv \emptyset\}$ and $\mathscr{C} = \emptyset$, hence $\theta(X) = \mathtt{End}$, $\theta(x) = \emptyset$ and the thesis follows by Axiom [ι-End].

Axiom [A-Cycle] We have $\mathscr{E} = \{X \Equiv Y\}$, $E = \{x \Equiv y\}$, $\mathscr{C} = \emptyset$, and $\mathscr{S} = \mathscr{S}', (Y, \mathbb{M}, y)$. Then, $\theta(X) = \theta(Y)$, $\theta(x) = \theta(y)$ and the thesis follows by Axiom [ι-Cycle].

Rule [A-Comm] We have $\mathbb{M} \equiv \mathsf{p}[\mathsf{q}!\{\lambda_i.P_i\}_{i \in I}] \parallel \mathsf{q}[\mathsf{p}?\{\lambda_j.Q_j\}_{j \in J}] \parallel \mathbb{M}'$ with $I \subseteq J$ and $\mathscr{S}, (X, \mathbb{M}, x) \vdash (Y_i, \mathbb{M}_i, y_i) \Rightarrow (\mathscr{E}_i, E_i, \mathscr{C}_i)$ with Y_i, y_i fresh and $\mathbb{M}_i \equiv \mathsf{p}[P_i] \parallel \mathsf{q}[Q_i] \parallel \mathbb{M}'$ and $\mathscr{E} = \{X \Equiv \mathsf{p} \to \mathsf{q} : \{\lambda_i.Y_i\}_{i \in I}\} \cup \bigcup_{i \in I} \mathscr{E}_i$, $E = \{x \Equiv \bigcup_{i \in I} y_i\} \cup \bigcup_{i \in I} E_i$, $\mathscr{C} = \{(\mathsf{prt}(Y_i) \cup y_i) \setminus \{\mathsf{p}, \mathsf{q}\} \doteq \mathsf{prt}(\mathbb{M}') \mid \forall i \in I\} \cup \bigcup_{i \in I} \mathscr{C}_i$. Since $\mathscr{E}_i \subseteq \mathscr{E}, E_i \subseteq E$, and $\mathscr{C}_i \subseteq \mathscr{C}$, we have $\theta \in \mathsf{sol}(\mathscr{E}_i, E_i, \mathscr{C}_i)$ for all $i \in I$. Being $\theta \in \mathsf{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$, Lemma 4.3 implies $\mathsf{prt}(\theta(X)) \cup \theta(x) = \mathsf{prt}(\mathbb{M})$. So we get that θ agrees with $\mathscr{S}, (X, \mathbb{M}, x)$. Then, by the induction hypothesis, we have $\mathscr{S}\theta, (\theta(X), \mathbb{M}, \theta(x)) \vdash^\iota_{\theta(y_i)} \mathbb{M}_i : \theta(Y_i)$ for all $i \in I$. The thesis follows by Rule [ι-Comm], since $\theta(X) = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\theta(Y_i)\}_{i \in I}$ and $\theta(x) = \bigcup_{i \in I} \theta(y_i)$.

Rule [A-Weak] We have $\mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $\mathscr{P} = \mathsf{prt}(\mathbb{M}_2) \neq \emptyset$ and $\mathscr{S}, (X, \mathbb{M}_1 \parallel \mathbb{M}_2, x) \vdash (Y, \mathbb{M}_1, y) \Rightarrow (\mathscr{E}_1, E_1, \mathscr{C})$ and $\mathscr{E} = \{X \Equiv Y\} \cup \mathscr{E}_1$ and $E = \{x \Equiv y \cup \mathscr{P}\} \cup E_1$. Being $\theta \in \mathsf{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$, Lemma 4.3

implies $\mathrm{prt}(\theta(X)) \cup \theta(x) = \mathrm{prt}(\mathbb{M})$. So we get that θ agrees with $\mathscr{S}, (X, \mathbb{M}, x)$. Then, by the induction hypothesis, we have $\mathscr{S}\theta, (\theta(X), \mathbb{M}, \theta(x)) \vdash^{\iota}_{\theta(y)} \mathbb{M}_1 : \theta(Y)$. The thesis follows by Rule [ι-Weak]. $\qquad\square$

Lemma 4.5 *If $\mathscr{N} \vdash^{\iota}_{\mathscr{P}} \mathbb{M} : \mathsf{G}$ and $\mathrm{prt}(\mathsf{G}') \cup \mathscr{P}' = \mathrm{prt}(\mathbb{M}')$ for all $(\mathsf{G}', \mathbb{M}', \mathscr{P}') \in \mathscr{N}$, then, for all $\mathscr{S}$, X, x and σ such that $X, x \notin \mathrm{vars}(\mathscr{S})$, $\mathrm{dom}(\sigma) = \mathrm{vars}(\mathscr{S})$ and $\mathscr{S}\sigma = \mathscr{N}$, there are $\mathscr{E}$, E, $\mathscr{C}$ and θ such that $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ and $\theta \in \mathrm{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$ and $\mathrm{dom}(\theta) = \mathrm{vars}(\mathscr{E}) \cup \mathrm{vars}(E) \cup \mathrm{vars}(\mathscr{S})$ and $\sigma \preceq \theta$ and $\theta(X) = \mathsf{G}$ and $\theta(x) = \mathscr{P}$.*

Proof. By induction on the derivation of $\mathscr{N} \vdash^{\iota}_{\mathscr{P}} \mathbb{M} : \mathsf{G}$. It is easy to verify that $\mathscr{N} \vdash^{\iota}_{\mathscr{P}} \mathbb{M} : \mathsf{G}$ implies $\mathrm{prt}(\mathsf{G}) \cup \mathscr{P} = \mathrm{prt}(\mathbb{M})$.

Axiom [ι-End] The thesis is immediate by Axiom [A-End] taking $\theta = \sigma + \{X \mapsto \mathtt{End}, x \mapsto \emptyset\}$.

Axiom [ι-Cycle] In this case we have $\mathscr{N} = \mathscr{N}', (\mathsf{G}, \mathbb{M}, \mathscr{P})$, then $\mathscr{S} = \mathscr{S}', (Y, \mathbb{M}, y)$ and $\sigma(Y) = \mathsf{G}$ and $\sigma(y) = \mathscr{P}$. By Axiom [A-Cycle], we get $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\{X \doteq Y\}, \{x \doteq y\}, \emptyset)$, hence $\theta = \sigma + \{X \mapsto \mathsf{G}, x \mapsto \mathscr{P}\}$ is a solution of $\{X \doteq Y\}$ and of $\{x \doteq y\}$, which agrees with $\mathscr{S}$, being $\mathrm{prt}(\mathsf{G}) \cup \mathscr{P} = \mathrm{prt}(\mathbb{M})$ as needed.

Rule [ι-Comm] In this case we have $\mathbb{M} \equiv \mathsf{p}[\mathsf{q}!\{\lambda_i.P_i\}_{i \in I}] \parallel \mathsf{q}[\mathsf{p}?\{\lambda_j.Q_j\}_{j \in J}] \parallel \mathbb{M}'$ with $I \subseteq J$ and $\mathsf{G} = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\mathsf{G}_i\}_{i \in I}$ and $\mathscr{N}, (\mathsf{G}, \mathbb{M}, \mathscr{P}) \vdash^{\iota}_{\mathscr{P}_i} \mathbb{M}_i : \mathsf{G}_i$ with $\mathbb{M}_i \equiv \mathsf{p}[P_i] \parallel \mathsf{q}[Q_i] \parallel \mathbb{M}'$ and $(\mathrm{prt}(\mathsf{G}_i) \cup \mathscr{P}_i) \setminus \{\mathsf{p}, \mathsf{q}\} = \mathrm{prt}(\mathbb{M}')$, for all $i \in I$. This last condition implies $\mathrm{prt}(\mathsf{G}) \cup \mathscr{P} = \mathrm{prt}(\mathbb{M})$. Set $\sigma' = \sigma + \{X \mapsto \mathsf{G}, x \mapsto \mathscr{P}\}$ and $\mathscr{S}' = \mathscr{S}, (X, \mathbb{M}, x)$, then, by the induction hypothesis, we get that there are $\mathscr{E}_i, E_i, \mathscr{C}_i$ and θ_i such that $\mathscr{S}' \vdash (Y_i, \mathbb{M}_i, y_i) \Rightarrow (\mathscr{E}_i, E_i, \mathscr{C}_i)$ and $\theta_i \in \mathrm{sol}_{\mathscr{S}'}(\mathscr{E}_i, E_i, \mathscr{C}_i)$ and $\mathrm{dom}(\theta_i) = \mathrm{vars}(\mathscr{E}_i) \cup \mathrm{vars}(E_i) \cup \mathrm{vars}(\mathscr{S}')$ and $\sigma' \preceq \theta_i$ and $\theta_i(Y_i) = \mathsf{G}_i$ and $\theta_i(y_i) = \mathscr{P}_i$, for all $i \in I$. We can assume that $j \neq l$ implies $Y_j \neq Y_l$ and $\mathrm{dom}(\mathscr{E}_j) \cap \mathrm{dom}(\mathscr{E}_l) = \emptyset$ and $y_j \neq y_l$ and $\mathrm{dom}(E_j) \cap \mathrm{dom}(E_l) = \emptyset$ for all $j, l \in I$, because the algorithm always introduces fresh variables. This implies $\mathrm{dom}(\theta_j) \cap \mathrm{dom}(\theta_l) = \{X, x\}$ for all $j \neq l$, and so $\theta = \sum_{i \in I} \theta_i$ is well defined. Moreover, we have $\theta \in \mathrm{sol}_{\mathscr{S}'}(\mathscr{E}_i, E_i, \mathscr{C}_i)$ and $\sigma \preceq \theta$ and $\theta(X) = \mathsf{G}$ and $\theta(x) = \mathscr{P}$, as $\sigma \preceq \sigma'$ and $\sigma' \preceq \theta_i \preceq \theta$ for all $i \in I$. From $(\mathrm{prt}(\mathsf{G}_i) \cup \mathscr{P}_i) \setminus \{\mathsf{p}, \mathsf{q}\} = \mathrm{prt}(\mathbb{M}')$ we get $(\mathrm{prt}(\theta(Y_i)) \cup \theta(y_i)) \setminus \{\mathsf{p}, \mathsf{q}\} = \mathrm{prt}(\mathbb{M}')$ for all $i \in I$. By Rule [A-Comm] we get $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ with $\mathscr{E} = \{X \doteq \mathsf{p} \to \mathsf{q} : \{\lambda_i.Y_i\}_{i \in I}\} \cup \bigcup_{i \in I} \mathscr{E}_i$ and $E = \{x \doteq \cup_{i \in I} y_i\} \cup \bigcup_{i \in I} E_i$ and $\mathscr{C} = \{(\mathrm{prt}(Y_i) \cup y_i) \setminus \{\mathsf{p}, \mathsf{q}\} \doteq \mathrm{prt}(\mathbb{M}) \mid \forall i \in I\} \cup \bigcup_{i \in I} \mathscr{C}_i$ and $\theta \in \mathrm{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C})$, since $\theta(X) = \mathsf{G} = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\mathsf{G}_i\}_{i \in I} = \mathsf{p} \to \mathsf{q} : \{\lambda_i.\theta_i(Y_i)\}_{i \in I} = (\mathsf{p} \to \mathsf{q} : \{\lambda_i.Y_i\}_{i \in I})\theta$ and $\theta(x) = \mathscr{P} = \cup_{i \in I} \mathscr{P}_i = \cup_{i \in I} \theta(y_i) = (\cup_{i \in I} y_i)\theta$ and $\sigma \preceq \theta$.

Rule [ι-Weak] We have $\mathscr{N}, (\mathsf{G}, \mathbb{M}_1 \parallel \mathbb{M}_2, \mathscr{P}_1 \cup \mathscr{P}_2) \vdash^{\iota}_{\mathscr{P}_1} \mathbb{M}_1 : \mathsf{G}$ and $\mathscr{P}_2 = \mathrm{prt}(\mathbb{M}_2) \neq \emptyset$ and $\mathrm{prt}(\mathsf{G}) \cup \mathscr{P} = \mathrm{prt}(\mathbb{M})$, where $\mathbb{M} \equiv \mathbb{M}_1 \parallel \mathbb{M}_2$ and $\mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2$. Set $\sigma' = \sigma + \{X \mapsto \mathsf{G}, x \mapsto \mathscr{P}_1 \cup \mathscr{P}_2\}$ and $\mathscr{S}' = \mathscr{S}, (X, \mathbb{M}, x)$, then, by the induction hypothesis, we get that there are $\mathscr{E}_1, E_1, \mathscr{C}_1$ and θ such that $\mathscr{S}' \vdash (Y, \mathbb{M}_1, y) \Rightarrow (\mathscr{E}_1, E_1, \mathscr{C}_1)$ and $\theta \in \mathrm{sol}_{\mathscr{S}'}(\mathscr{E}_1, E_1, \mathscr{C}_1)$ and $\mathrm{dom}(\theta) = \mathrm{vars}(\mathscr{E}_1) \cup \mathrm{vars}(E_1) \cup \mathrm{vars}(\mathscr{S}')$ and $\sigma' \preceq \theta$ and $\theta(Y) = \mathsf{G}$ and $\theta(y) = \mathscr{P}_1$. By Rule [A-Weak] we get $\mathscr{S} \vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C}_1)$ with $\mathscr{E} = \{X \doteq Y\} \cup \mathscr{E}_1$ and $E = \{x \doteq y \cup \mathscr{P}_2\} \cup E_1$ and $\theta \in \mathrm{sol}_{\mathscr{S}}(\mathscr{E}, E, \mathscr{C}_1)$, since $\theta(X) = \mathsf{G} = \theta(Y)$ and $\theta(x) = \mathscr{P} = \mathscr{P}_1 \cup \mathscr{P}_2 = \theta(y) \cup \mathscr{P}_2 = (y \cup \mathscr{P}_2)\theta$ and $\sigma \preceq \theta$. $\qquad\square$

Theorem 4.6 (Soundness and completeness of inference)

1. *If $\vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$, then $\theta(X) \vdash_{\theta(x)} \mathbb{M}$ for all $\theta \in \mathrm{sol}(\mathscr{E}, E, \mathscr{C})$.*

2. *If $\mathsf{G} \vdash_{\mathscr{P}} \mathbb{M}$, then there are $\mathscr{E}$, E, $\mathscr{C}$ and θ such that $\vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ and $\theta \in \mathrm{sol}(\mathscr{E}, E, \mathscr{C})$ and $\theta(X) = \mathsf{G}$ and $\theta(x) = \mathscr{P}$.*

Proof. (1). By Lemma 4.4 $\vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ implies $\vdash^{\iota}_{\theta(x)} \mathbb{M} : \theta(X)$ for all $\theta \in \mathrm{sol}(\mathscr{E}, E, \mathscr{C})$. This is enough, since $\vdash^{\iota}_{\theta(x)} \mathbb{M} : \theta(X)$ gives $\theta(X) \vdash_{\theta(x)} \mathbb{M}$.

(2). From $\mathsf{G} \vdash_{\mathscr{P}} \mathbb{M}$ we get $\vdash^{\iota}_{\mathscr{P}} \mathbb{M} : \mathsf{G}$. By Lemma 4.5 this implies that there are $\mathscr{E}$, E, $\mathscr{C}$ and θ such that $\vdash (X, \mathbb{M}, x) \Rightarrow (\mathscr{E}, E, \mathscr{C})$ and $\theta \in \mathrm{sol}(\mathscr{E}, E, \mathscr{C})$ and $\theta(X) = \mathsf{G}$ and $\theta(x) = \mathscr{P}$. $\qquad\square$

Remark 4.7 (Termination) As happens for (co)SLD-resolution in logic programming, the termination of the inference algorithm depends on the choice of a resolution strategy. Indeed, we have many sources of non-determinism: we have to select two participants of the session with matching processes and expand them using Rule [A-COMM], or ignore part of the session using Rule [A-WEAK] or try to close a cycle using the Axiom [A-CYCLE]. A standard way to obtain a sound and complete resolution strategy is to build a tree where all such choices are performed in parallel and then visit the tree using a breadth-first strategy. The tree is potentially infinite in depth, but it is finitely branching, since at each point we have only finitely many different choices, hence this strategy necessarily enumerates all solutions.

Remark 4.8 (Use of Rule [A-WEAK]) Note that in a $\vdash^{\iota}_{\mathscr{P}}$ derivation the triple in the premise of Rule [ι-WEAK] can never be used in an application of Axiom [ι-CYCLE]. This – as already hinted at in Example 3.2 – immediately implies that Rule [WEAK] is not strictly necessary inside infinite branches of $\vdash_{\mathscr{P}}$ derivations. Moreover, a slight simplification of the algorithm can be got since, in the step corresponding to Rule [A-WEAK], we could avoid adding the goal $(X, \mathbb{M}_1 \parallel \mathbb{M}_2, x)$ to the current set of goals. This would reduce the number of goals to be checked during the step corresponding to Axiom [ι-CYCLE]. Rule [A-WEAK] turns out to be necessary, instead, when applying the algorithm to sessions where non-ignored participants expose a finite behaviour, like $\mathsf{p}[P_2] \parallel \mathsf{q}[Q_2] \parallel \mathsf{u}[U]$ in Example 3.2. Also the typing of stuck sessions with recursive processes like $\mathsf{p}[P] \parallel \mathsf{q}[Q]$ where $P = \mathsf{q}!\lambda.P$ and $Q = \mathsf{p}!\lambda.Q$ requires the use of Rule [A-WEAK]. $\diamond$

5   Concluding Remarks, Related and Future Works

Lock-freedom is definitely a relevant and widely investigated communication property of concurrent and distributed systems. It ensures absence of *locks*, a lock being a reachable configuration where a communication action of a participant remains pending in any possible continuation of the system. In case the participant prevented to progress be p, such configuration is called a p-*lock* (see [5] for an abstract definition of Lock-freedom). Lock-freedom corresponds to the notion of liveness in [20, 22] where the synchronous communication is channel-based. Sometimes properties different from what we intend are referred to by "Lock-freedom": for instance the notion of Lock-freedom in [19], under fair scheduling, corresponds to what [28] and [5] refer to as *strong Lock-freedom*.

Various formalisms and methodologies have been developed in order to prove Lock-freedom while others do ensure Lock-freedom by construction. Among the former there are type assignment systems where typability entails Lock-freedom, both for asynchronous [26] and synchronous [3] communications.

Lock-freedom is quite a strong property: it entails Deadlock-freedom, whereas the vice versa does not hold. In several actual scenarios, lighter forms of Lock-freedom would however suffice. For instance in clients/servers scenarios where one can accept some servers to get locked after their interactions with the clients have been completed.

In the present paper we developed a type assignment system where typability ensures $\mathscr{P}$-*excluded Lock-freedom*: the absence of p-locks for each participant p not belonging to $\mathscr{P}$. This is achieved by means of "partial" typability, i.e. by disregarding typability of (sub)processes of participants that we can safely assume to get possibly locked. Multiparty sessions (parallel compositions of named processes) are (partially) typed by *global types*, which in turn describe the overall interactions inside the multiparty sessions. Our partial typability ensures also that the behaviours of the non-ignored participants adhere to what the global type describes. As far as we know, there are not other formalisms dealing with properties like $\mathscr{P}$-excluded Lock-freedom.

Our partial typing is reminiscent of connecting communications, a notion introduced in [18] and further investigated in [8, 10] in order to describe protocols with optional participants. The intuition behind connecting communications is that in some parts of the protocol, delimited by a choice construct, some participants may be optional, namely they are "invited" to join the interaction only in some branches of the choice, by means of connecting communications. As argued in [18, 8, 10], this feature allows for a more natural description of typical communication protocols. In [10], connecting communications also enable to express conditional delegation: this will be obtained by writing a choice where the delegation appears only in some branches of the choice, following a connecting communication. The participants offering connecting communications should be ignored in the present type system. An advantage of connecting communications over partial typing is that only participants offering connecting inputs can be stuck. The disadvantage is that the typing rules are more requiring, so many interesting sessions can be partially typed but cannot be typed by means of connecting communications.

In designing type inference we took inspiration from [13], where inputs and outputs are split in global types in order to better describe asynchronous communication. Our inference algorithm is related as goal, but very different as methodology, to the algorithm in [23], which builds global graphs from sets of communicating finite state machines satisfying suitable conditions. We are planning to implement our type inference algorithm.

Unlike many MPST formalisms in the literature, like [17], we type sessions with global types without recurring to local types and projections. It would be interesting to investigate the possibility of extending the standard projection operator to a relation between global types and possibly non lock-free local behaviours. Other simplifications of our calculus are the absence of values in messages and the unicity of channels. While we can easily enrich messages with values, allowing more than one channel requires sophisticated type systems in order to get Lock-freedom [26].

The following example shows a further direction for investigation of partial typing, namely to describe and analyse privacy matters.

Example [**Partial typing for privacy**] The communications written in global types can be viewed as public, while the others can be viewed as private. For example Alice and Bob want to discuss privately which version of a game would be the most suitable for their son Carl, who asked for it as birthday present. Taking participants a, b and c to incarnate, respectively, Alice, Bob and Carl this scenario can be represented by the session

$$\mathbb{M} \equiv \mathsf{a}[\mathsf{c}?\textsc{present}.P] \parallel \mathsf{b}[\mathsf{c}?\textsc{present}.Q] \parallel \mathsf{c}[\mathsf{a}!\textsc{present}.\mathsf{b}!\textsc{present}]$$

where $P = \mathsf{b}!\{\textsc{bla}.\mathsf{b}?\textsc{bla'}.P, \textsc{ok}\}$ and $Q = \mathsf{a}?\{\textsc{bla}.\mathsf{a}!\textsc{bla'}.P, \textsc{ok}\}$.
A suitable global type is $\mathsf{G} = \mathsf{c} \to \mathsf{a}{:}\textsc{present}.\mathsf{c} \to \mathsf{b}{:}\textsc{present}$. We can in fact derive $\mathsf{G} \vdash_{\{\mathsf{a},\mathsf{b}\}} \mathbb{M}$.   ◇

We also plan to investigate partial typing for asynchronous communications, possibly modifying the type system of [13]. An advantage of that type system is the possibility of anticipating outputs over inputs without requiring the asynchronous subtyping of [25], which is known to be undecidable [6, 24]. A difficulty will come from the larger freedom in choosing the order of interactions due to the splitting between writing and reading messages on a queue.

References

[1] Jirí Adámek, Stefan Milius & Jiri Velebil (2006): *Iterative algebras at work*. *Mathematical Structures in Computer Science* 16(6), pp. 1085–1131, doi:10.1017/S0960129506005706.

[2] Davide Ancona & Agostino Dovier (2015): *A theoretical perspective of coinductive logic programming*. *Fundamenta Informaticae* 140(3-4), pp. 221–246, doi:10.3233/FI-2015-1252.

[3] Franco Barbanera, Mariangiola Dezani-Ciancaglini & Ugo de'Liguoro (2022): *Open compliance in multiparty sessions*. In S. Lizeth Tapia Tarifa & José Proença, editors: *FACS*, *LNCS* 13712, Springer, pp. 222–243, doi:10.1007/978-3-031-20872-0_13.

[4] Franco Barbanera, Mariangiola Dezani-Ciancaglini, Ivan Lanese & Emilio Tuosto (2021): *Composition and decomposition of multiparty sessions*. *Journal of Logic and Algebraic Methods in Programming* 119, p. 100620, doi:10.1016/j.jlamp.2020.100620.

[5] Franco Barbanera, Ivan Lanese & Emilio Tuosto (2022): *Formal choreographic languages*. In Maurice H. ter Beek & Marjan Sirjani, editors: *COORDINATION*, *LNCS* 13271, Springer, pp. 121–139, doi:10.1007/978-3-031-08143-9_8.

[6] Mario Bravetti, Marco Carbone & Gianluigi Zavattaro (2017): *Undecidability of asynchronous session subtyping*. *Information and Computation* 256, pp. 300–320, doi:10.1016/j.ic.2017.07.010.

[7] Giuseppe Castagna, Nils Gesbert & Luca Padovani (2009): *A theory of contracts for Web services*. *ACM Transaction on Programming Languages and Systems* 31(5), pp. 19:1–19:61, doi:10.1145/1538917.1538920.

[8] Ilaria Castellani, Mariangiola Dezani-Ciancaglini & Paola Giannini (2019): *Reversible sessions with flexible choices*. *Acta Informatica* 56(7), pp. 553–583, doi:10.1007/s00236-019-00332-y.

[9] Ilaria Castellani, Mariangiola Dezani-Ciancaglini & Paola Giannini (2022): *Asynchronous sessions with input races*. In Marco Carbone & Rumyana Neykova, editors: *PLACES*, *EPTCS* 356, Open Publishing Association, pp. 12–23, doi:10.4204/EPTCS.356.2.

[10] Ilaria Castellani, Mariangiola Dezani-Ciancaglini, Paola Giannini & Ross Horne (2020): *Global types with internal delegation*. *Theoretical Computer Science* 807, pp. 128–153, doi:10.1016/j.tcs.2019.09.027.

[11] Luca Ciccone, Francesco Dagnino & Luca Padovani (2022): *Fair termination of multiparty sessions*. In Karim Ali & Jan Vitek, editors: *ECOOP*, *LIPIcs* 222, Schloss Dagstuhl - Leibniz-Zentrum für Informatik, pp. 26:1–26:26, doi:10.4230/LIPIcs.ECOOP.2022.26.

[12] Bruno Courcelle (1983): *Fundamental properties of infinite trees*. *Theoretical Computer Science* 25, pp. 95–169, doi:10.1016/0304-3975(83)90059-2.

[13] Francesco Dagnino, Paola Giannini & Mariangiola Dezani-Ciancaglini (2023): *Deconfined global types for asynchronous sessions*. *Logical Methods in Computer Science* 19(1), pp. 1–41, doi:10.46298/lmcs-19(1:3)2023.

[14] Romain Demangeon & Kohei Honda (2012): *Nested protocols in session types*. In Maciej Koutny & Irek Ulidowski, editors: *CONCUR*, *LNCS* 7454, Springer, pp. 272–286, doi:10.1007/978-3-642-32940-1_20.

[15] Rob van Glabbeek, Peter Höfner & Ross Horne (2021): *Assuming just enough fairness to make session types complete for lock-freedom*. In Leonid Libkin, editor: *LICS*, ACM Press, pp. 1–13, doi:10.1109/LICS52264.2021.9470531.

[16] Kohei Honda, Nobuko Yoshida & Marco Carbone (2008): *Multiparty asynchronous session types*. In George C. Necula & Philip Wadler, editors: *POPL*, ACM Press, pp. 273–284, doi:10.1145/1328897.1328472.

[17] Kohei Honda, Nobuko Yoshida & Marco Carbone (2016): *Multiparty asynchronous session types*. *Journal of the ACM* 63(1), pp. 9:1–9:67, doi:10.1145/2827695.

[18] Raymond Hu & Nobuko Yoshida (2017): *Explicit connection actions in multiparty session types*. In: *FASE*, *LNCS* 10202, Springer, pp. 116–133, doi:10.1007/978-3-662-54494-5.

[19] Naoki Kobayashi (2002): *A type system for lock-free processes*. *Information and Computation* 177(2), pp. 122–159, doi:10.1006/inco.2002.3171.

[20] Naoki Kobayashi & Davide Sangiorgi (2010): *A hybrid type system for lock-freedom of mobile processes*. ACM Transactions on Programming Languages and Systems 32(5), pp. 16:1–16:49, doi:10.1145/1745312.1745313.

[21] Dexter Kozen & Alexandra Silva (2017): *Practical Coinduction*. Mathematical Structures in Computer Science 27(7), pp. 1132–1152, doi:10.1017/S0960129515000493.

[22] Julien Lange, Nicholas Ng, Bernardo Toninho & Nobuko Yoshida (2017): *Fencing off Go: liveness and safety for channel-based programming*. In Giuseppe Castagna & Andrew D. Gordon, editors: *POPL*, ACM Press, pp. 748–761, doi:10.1145/3009837.3009847.

[23] Julien Lange, Emilio Tuosto & Nobuko Yoshida (2015): *From communicating machines to graphical choreographies*. In Sriram K. Rajamani & David Walker, editors: *POPL*, ACM Press, pp. 221–232, doi:10.1145/2676726.2676964.

[24] Julien Lange & Nobuko Yoshida (2017): *On the undecidability of asynchronous session subtyping*. In Javier Esparza & Andrzej S. Murawski, editors: *FOSSACS*, *LNCS* 10203, Springer, pp. 441–457, doi:10.1007/978-3-662-54458-7_26.

[25] Dimitris Mostrous, Nobuko Yoshida & Kohei Honda (2009): *Global principal typing in partially commutative asynchronous sessions*. In Giuseppe Castagna, editor: *ESOP*, *LNCS* 5502, Springer, pp. 316–332, doi:10.1007/978-3-642-00590-9_23.

[26] Luca Padovani (2014): *Deadlock and lock freedom in the linear π-calculus*. In Thomas A. Henzinger & Dale Miller, editors: *CSL-LICS*, ACM Press, pp. 72:1–72:10, doi:10.1007/978-3-662-43376-8_10.

[27] Benjamin C. Pierce (2002): *Types and Programming Languages*. MIT Press.

[28] Paula Severi & Mariangiola Dezani-Ciancaglini (2019): *Observational equivalence for multiparty sessions*. Fundamenta Informaticae 167, pp. 267–305, doi:10.3233/FI-2019-1863.

[29] Luke Simon (2006): *Extending logic programming with coinduction*. Ph.D. thesis, University of Texas at Dallas.

[30] Luke Simon, Ajay Bansal, Ajay Mallya & Gopal Gupta (2007): *Co-logic programming: extending logic programming with coinduction*. In Lars Arge, Christian Cachin, Tomasz Jurdzinski & Andrzej Tarlecki, editors: *ICALP*, *LNCS* 4596, Springer, pp. 472–483, doi:10.1007/11799573_25.

[31] Luke Simon, Ajay Mallya, Ajay Bansal & Gopal Gupta (2006): *Coinductive logic programming*. In Sandro Etalle & Miroslaw Truszczynski, editors: *ICLP*, *LNCS* 4079, Springer, pp. 330–345, doi:10.1007/11799573_25.

Comprehensive Specification and Formal Analysis of Attestation Mechanisms in Confidential Computing

Muhammad Usama Sardar[1], Thomas Fossati[2] and Simon Frost[2]

[1]TU Dresden
[2]Arm Ltd.

Confidential Computing (CC) using hardware-based Trusted Execution Environments (TEEs) has emerged as a promising solution for protecting sensitive data in all forms. One of the fundamental characteristics of such TEEs is remote attestation [4] which provides mechanisms for securely measuring and reporting the state of the remote platform and computing environment to a user. We present a novel approach combining TEE-agnostic attestation architecture and formal analysis enabling comprehensive and rigorous security analysis of attestation mechanisms in CC. We demonstrate the application of our approach for three prominent industrial representatives, namely Arm Confidential Compute Architecture (CCA) [1, 2] in architecture lead solutions, Intel Trust Domain Extensions (TDX) [5] in vendor solutions, and Secure CONtainer Environment (SCONE) [3] in frameworks. For each of these solutions, we provide a comprehensive specification of all phases of the attestation mechanism in confidential computing, namely provisioning, initialization, and attestation protocol. Our approach reveals design and security issues in Intel TDX and SCONE attestation. The work is currently under submission at another venue with formal proceedings [6].

References

[1] Arm Ltd.: Arm Realm Management Extension (RME) System Architecture. Tech. rep. (2021-2022), https://developer.arm.com/documentation/den0129

[2] Arm Ltd.: Realm Management Monitor specification. Tech. rep. (2021-2022), https://developer.arm.com/documentation/den0137

[3] Arnautov, S., Trach, B., Gregor, F., Knauth, T., Martin, A., Priebe, C., Lind, J., Muthukumaran, D., O'keeffe, D., Stillwell, M.L., et al.: SCONE: Secure linux containers with Intel SGX. In: USENIX Symposium on Operating Systems Design and Implementation. pp. 689–703 (2016), https://www.usenix.org/conference/osdi16/technical-sessions/presentation/arnautov

[4] Guttman, J.D., Ramsdell, J.D.: Understanding Attestation: Analyzing Protocols that Use Quotes. In: Security and Trust Management, pp. 89–106 (2019), http://link.springer.com/10.1007/978-3-030-31511-5_6

[5] Intel: Intel ® Trust Domain Extensions (aug 2021), https://cdrdv2.intel.com/v1/dl/getContent/690419

[6] Sardar, M.U., Fossati, T., Frost, S.: SoK: Attestation in Confidential Computing (2023), https://www.researchgate.net/publication/367284929_SoK_Attestation_in_Confidential_Computing